

Ruprecht-Karls-Universität Heidelberg
Institut für Informatik
Lehrstuhl für Parallele und Verteilte Systeme

Masterarbeit

Design and Implementation of a tool for
automatic, non-trivial code guideline
checking Automated detection of violations
of clean code Automated Checking of Clean
Code Guidelines for Python «—- jedes Wort
was nicht and a ist capitalisieren

Name: Enrico Kaack
Matrikelnummer: 3534472
Betreuer: Name des Betreuers
Datum der Abgabe: dd.mm.yyyy

Ich versichere, dass ich diese Master-Arbeit selbstständig verfasst und nur die angegebenen Quellen und Hilfsmittel verwendet habe.

Heidelberg, dd.mm.yyyy

Zusammenfassung

Abstract auf Deutsch

Abstract

This is the abstract.

Contents

List of Figures	1
List of Tables	2
1 Introduction	3
1.1 Motivation	3
1.2 Goals	3
1.3 Structure	3
2 Background and Related Work	4
2.1 Code Quality	4
2.2 Clean Code	5
2.2.1 General Rules	6
2.2.2 Naming	6
2.2.3 Functions	7
2.2.4 Comments	9
2.2.5 Data Structures and Objects	10
2.2.6 Classes	11
2.2.7 Exception Handling	11
2.2.8 Code Smells	12
2.2.9 Additional Guidelines	12
2.3 Quantitative Metrics for Code Quality	13
2.3.1 Cyclomatic Complexity	13
2.3.2 Halstead complexity measures	14
2.3.3 Software Maintainability Index	15
2.4 Tools for Code Quality Analysis	16
2.4.1 PyLint	17
2.4.2 PMD Source Code Analyzer Project	17
2.4.3 Codacy	18

2.4.4	Sonarqube	19
3	Approach	22
3.1	Clean Code Analysis Plattform	22
3.1.1	Architecture	23
3.1.2	Extensions	32
3.2	Clean Code Classification	34
3.2.1	Challenges	34
3.2.2	Dataset	37
3.2.3	Processing	40
3.2.4	Models	45
4	Quantitative Evaluation	48
4.1	Research Questions	48
4.1.1	RQ1: Can the Clean Code Analysis Platform be a useful addition to developers workflow besides existing tools?	48
4.1.2	RQ2: What models perform well on detecting non-clean code? . .	50
5	Conclusion	56
	Bibliography	57

List of Figures

3.1	Output of the HTML Output Plugin, displayed in a browser.	33
3.2	Example commit messages that underline the inconsistency in commit message naming	35
3.3	Split of the dataset. First, 20% of all files are seperated into a holdout set for later testing. The remaining 80% of the files are split into 90% training and 10% test data.	40
3.4	Schematic representation of the pipeline for RQ2 with configurable parameters.	41
3.5	Pipeline for code manipulation and evaluation. TODO detailed description	44

List of Tables

3.1	Class distribution for the train/test dataset before the split. Label [0] represents clean code and label [1] marks problematic code of the corresponding category.	36
3.2	TODO lstm	38
3.3	General and problem specific metrics for the train, test and holdout set. The lines of code containing a problem and the problems per file are similiar.	39
3.4	Label frequency for all problem types on the train, test and holdout set. .	39
4.1	Best performing classifier for RQ2 (based on f1 score).	52
A1	random forest	58
A2	gradient boosting classifier	59
A3	lstm	60

1 Introduction

This chapter contains an overview of the topic as well as the goals and contributions of your work. Description of the domain Description of the problem Summary of the approach Outline of own contributions and results

1.1 Motivation

1.2 Goals

1.3 Structure

Here you describe the structure of the thesis. For example:

In Kapitel 2 werden grundlegende Methoden für diese Arbeit vorgestellt.

2 Background and Related Work

General description of the relevant methods/techniques

What has been done so far to address the problem (closer related work)

Possible weaknesses of the existing approaches

2.1 Code Quality

Code Quality describes the quality of source code with regard to understandability and readability. Developers can understand well-written code easily. High-quality code has an impact on onboarding new developers, writing new code and maintaining the existing code.

The onboarding of new developers is an investment of time and money in the developer. The faster a new developer can understand an existing code base, the faster the developer can start writing productive code and providing value.

Maintaining existing code and adding features is part of most software today (TODO source). Agile development is a methodology used in software development that reflects this requirement. Source Code is improved and changed with runnable software versions at the end of each iteration. From a business standpoint, the always-changing code is modeled by subscription-based contracts that include new features and bugfixes. The easiness to change source code is business-critical, and a high-quality code can affect this requirement in the following kinds [3]:

1. Well-written code makes it easy to determine the location and the way source code has to be changed.
2. A developer can implement changes more efficient in good code.
3. Easy to understand code can prevent unexpected side-effects and bugs when applying a change.
4. Changes can be validated easier.

The International Organization for Standardization provides the standard ISO/IEC 25000:2014 for “Systems and software Quality Requirements and Evaluation (SQuaRE)”[10].

Code Quality is measured by the following code characteristics:

1. Reliability
2. Performance efficiency
3. Security
4. Maintainability

Besides the mentioned maintainability characteristics, Code Quality also depends on reliability (like multi-threading and resource allocation handling), performance efficiency for efficient code execution, and security (like vulnerabilities to frequent attacks like SQL-injection).

2.2 Clean Code

Clean Code is a concept for high-quality code, coined by the book Clean Code by Robert C. Martin [7]. The root cause for unclean code is chaotic code. Developers produce chaotic code in a conflict between deadline pressure based on the visible output (the functionality of the software) and extra effort to make code more intuitive. The latter is not directly visible as productive output, although an accumulation of chaotic code reduces the productivity over time [7]. A bigger legacy system with chaotic code will slow down later modifications or additions of code. By following the Clean Code guidelines and best-practices, this productivity loss can be minimized.

The Clean Code techniques focus mainly on maintainability by providing intuitive code. This has a positive effect on the security and reliability aspect as well, since developers can find edge cases in non-logical behaviour more easily in intuitive code. Some of the following Clean Code principles may decrease performance efficiency. Still, in many software projects, developer performance is a more valuable resource than actual runtime performance (TODO source).

The following sections explain the clean code guidelines following the book by Robert C. Martin [7]. Since these rules are based on experience of the author, they are controversial. The critique will be explained in section TODO.

2.2.1 General Rules

Developers should follow the general rules consistently for developing new features or fixing bugs. They are the essential building blocks for the understandability and reliability of the code.

Follow standard conventions is the first rule. Programming languages have conventions on formatting, naming, etc (e.g. python¹). If developers follow these conventions, the code feels more familiar for other developers using the same conventions. It is also common practice for big open-source projects to have their own contributing guidelines with coding conventions. The Visual Studio Code GitHub repository contains a "How to contribute" documentation and a section on coding guidelines². Especially in such large open-source projects, many developers are working asynchronously on code. Therefore, having conventions and enforcing the compliance of the guidelines by rejecting pull requests prevents the code from becoming a mosaic of different coding styles.

The rule keep it simple is a summary of most rules in clean code. Simplicity in code makes it simpler and faster to understand for developers. A simpler software architecture enables a developer to modify code and checking all dependencies for potential side effects. Unnecessary complexity increases the time to understand and modify code and introduces bugs by having complex, non-obvious behaviors.

Everytime a developer touches the code, he should also improve the quality of the code or at least not worsen it. By doing a small fix and postponing the clean code principles, the code will become more chaotic until it is refactored. As a result, every change should keep at least keep the code level quality if not improving it.

2.2.2 Naming

It is important for understandable code to have good namings for variables, functions, types and classes. "Good" naming is an opinionated topic; The author describes the key components to good naming as follows: Names should be descriptive for the object. Abbreviations or mathematical annotations like *a1*, *a2* are not descriptive and do not provide information about the meaning. Implementations of mathematical expressions intentionally use the same terms as the expression itself. Since this increases the understandability for developers familiar with the mathematical expression, this can be seen as a valid exception. Descriptive names should include a verb or verb expression for

¹<https://www.python.org/dev/peps/pep-0008/>

²<https://github.com/Microsoft/vscode/wiki/Coding-Guidelines>

functions, since functions express an action. Conversely, class names should contain a noun to emphasize the object character of a class.

If the descriptive names are pronounceable too, it is easier to read and it is easier to talk about the code with other developers. Therefore, it is useful to make name longer but descriptive and pronounceable, especially since the autocomplete feature of IDEs will free the developer from typing the long name. Additionally, searching for long names works better than for short names, since long names are more likely to be unambiguous compared to shorter names. Short variables in a small scope (e.g. variable *i* in a loop) are not problematic, but using *i* in a large scope could be ambiguous and troublesome for searching.

An old naming convention is the Hungarian naming convention. In Hungarian naming, the type is encoded as a prefix of a variable name. Nowadays, this is seen as mostly redundant, since IDEs can infer and show the type automatically. The automatic type inference also prevents confusion for the reader if the type in the notation and the actual type are inconsistent. The same logic applies to a prefix for member variables and methods, since a IDE can automatically highlight those tokens.

2.2.3 Functions

Functions should be small in length. Exceeding 20 lines should not be necessary in most cases. If a function is small, it can be read more easily and without scrolling. Inside a function, if-, else- and while-statements should contain a function call for the condition and one function call for the body. By following the naming schemas in section 2.2.2, the function calls document the meaning of the condition in an intuitive way without the need for additional comments. Consequently, function calls replaces nested structures in a more readable way.

A general guideline for functions is to fulfill one purpose. This is a rather impractical view, since functions may have to call several functions subsequently to perform the required computation. Therefore, Robert C. Martin specifies that to one abstraction level per function. A low abstraction level handles data access and manipulation, e.g. string manipulation. On a middle abstraction layer, these low-level operations are orchestrated. On the next higher level, the mid-level functions are arranged etc. Switch-Statements violate the one purpose rule and is prone to be duplicated in several other code locations. Since they are a sometimes necessary construct, placing switch-statements into an Abstract Factory and creating different Subclasses for the different behaviours, the switch-statement can be replaced by polymorphism.

The number of function arguments should be three or lower. With many function arguments, it becomes harder to use the function. Additionally, testing becomes harder with more arguments; especially if all argument combinations should be tested. Functions with none or one argument simplifies testing alot. Since testing reduces the number of bugs, testable code is crucial. More arguments should be bundles in a config object or class, so many arguments are passed as one. Furthermore, using config objects enables grouping of arguments for the same context. A special focus lies on "output arguments"; these get passed as reference and are mutated in the function. Output arguments are common in low-level languages like C, but they require additional attention by the reader. Functions are expected to use the input values and return the output as a return value. Consequently, if a developer fails to notice an output argument, he may does not expect it to be mutated by the function and could face a logical error. Using function arguments as immutable is the most intuitive way. Although it is sometime necessary to mutate the input arguments in a function for performance reasons (TODO source).

A particular bad practice in the function body are side-effects. Side-effects happens, if a function modifies a variable outside its scope without explicitly mentioning this in the name. This leads to dependencies between the functions that are not obvious. A developer who uses the function checks the signature (meaning the name, input arguments and return type) and expects the function to do what the name suggests. If the function has a side effect, it is an unintended behaviour and will lead to mistakes that are time-consuming to debug. Especially side-effects that initialize other objects result in a time-dependency that is hard to identify and could be overseen at all.

Following the one purpose per function guideline, a function should either perform an action or retrieve information. Especially actions that return a value are unintuitive, since it is not clear if the return value encodes the success state or indicate something else. Furthermore, returning error codes violates this rule and errors should be indicated by raising an exception, if the language supports it. This allows a better separation between application logic and error handling code. The code for catching and handling an exception may be separated into an additional function, to provide a clean structure for higher level functions. In this case, error handling counts the one purpose of this function. A general, important principle for software engineering applies directly to functions: Don't repeat yourself. Repeated Code is dangerous and chaotic, since it requires changes in multiple locations if it has to be modified. This makes duplicated code very prone to copy and paste errors and small mistakes that may not be obvious

during development but will lead to a fatal crash at runtime. Many patterns and tools have been developed to mitigate duplicated code (TODO: some references to duplicated code detection etc).

2.2.4 Comments

Comments are part of every programming language and can play an important role in code quality. Good comments clarify the meaning of the code and help to understand the code. However, comments can become outdated and wrong or provide useless information. In a perfect world, the programming language and the programmer would be expressive enough that commenting is not required for clarity. Following the clean code guidelines for naming and functions makes many descriptive comments obsolete, since the function description is encoded in the function name.

Comments will not help to turn chaotic code into clean code. If a developer explains a line of code by a comment, it is often more helpful to call a function with a descriptive name to replace the comment. Since comments are not part of the program logic, if the explaining comment is not updated with the code, the comment gives misinformation to the reader like described before.

In brief, before writing a comment, the developer should think about expressing the same in code. Robert C. Martin gives some exceptions for good comments [7]:

Legal notes: Legal notes like copyright or license information and author mentions may be necessary. Although they should be short in link to an external licensing document in full extent.

Explaining comments: Some explaining comments can be helpful and are not easy to encode in normal code. For instance an explanation of special encodings or file formats are better to understand.

Intention: Explaining an intention in a comment that is not obvious by the source code is also a valid use for a comment and can help to understand code that otherwise would not make any sense.

Warning for consequences: Some parts of the code can have special consequences like not being thread safe or using many system resources. A warning can help a developer from using a function and having trouble.

Emphasize: A comment to emphasize a seemingly unimportant part of the code prevents breaking modifications of the code.

2.2.5 Data Structures and Objects

Data Structures and Objects can be easily mixed together, but their purpose is converse. Objects hide data behind abstractions and has functions that work with these data. Conversely, data structures expose the data and do not provide functionality. Based on this definition, the Law of Demeter (TODO source) is defined: Objects should know as little as possible about the implementation of objects it manipulates. As a result, code becomes decoupled, since objects do not depend on specific data structures of other objects. In object-oriented programming, a method *m* of an object *O* may only call methods of the following components:

- the object *O* itself
- the parameters of the method
- objects that are created within *m*'s scope
- instance variables of *O*

The method should not call methods of objects, that are returned from the allowed objects. If a method calls specific functions of these objects, it assumes this behaviour. As a result, the two objects are closely coupled. Such bad code practice is called train wreck, since subsequent method calls look like multiple train carriages. A split in multiple variables and method calls improves the Train Wreck. Nevertheless, method calls on objects assume knowledge of the internal implementation of the objects and would violate the Law of Demeter. However, if the methods are accessor functions, it would not violate the law, since the method calls are just an access to the data structure. Instead of calling a method of an object and calling a method on the returning object, the first called object could provide a function that takes care of calling the other object. This reduces the coupling between the three objects and would preserve the Law of Demeter in all objects.

Data structures are often used as Data Transfer Object (DTO) to communicate with other processes or services. These objects do not contain any functions, they only have accessible member variable. A specialisation of DTOs are Active Records, that contain additional methods for data storage, since they are used to represent a data source like a database. For both types, it is bad practice to insert business logic into these objects, since they should be treated as data structures. This creates a hybrid between an object and a data structure that should be avoided.

2.2.6 Classes

A Class should be small like functions. Instead of counting lines, the size of a class is the number of responsibilities. The name of the class should describe its responsibility, therefore, a single responsibility per class is intended. This is called Single-Responsibility-Principle and it restricts a class to one responsibility. A class should only have one reason for modification, consequently a class has a single responsibility (TODO source). The Single-Responsibility-Principle lead to a system of many small classes with one, clearly defined responsibility.

In addition, classes should have a high cohesion. Cohesion is high, if a method manipulates many instance variables. If all methods of a class manipulate all instance variables, the class is maximal cohesive. With a high cohesion, the methods and the class can be seen as a logical unit. A low cohesion indicates to split the class into multiple classes, since the methods are not a logical unit. As a result, classes gets smaller and will more likely comply to the Single-Responsibility-Principle. (TODO add metric for cohesion)

Dependencies of classes are critical, since they may be modified. If the class depends on specific details of the implementation, it has to be changed, if the dependency changes. Therefore, classes should depend on abstract classes that describe a concept, not the implementation details. Since the concept is unlikely to change if the implementation details change, the dependency is isolated.

2.2.7 Exception Handling

Section 2.2.3 briefly mentioned a clean code guideline for exception handling. Exceptions should be handled by raising and catching exceptions instead of error codes or flags. Error code handling introduces code, that is not coherent with main logic. Error codes have to be checked immediately and error handling has to be implemented directly. As a result, the two concerns main logic and error logic are mixed. By raising and catching errors, the error handling can be completely extracted into a separate function and can be removed from the main logic. Furthermore, the try-catch block enforces a transactional behaviour. At the end of either the try or the catch block, the application should be in a consistent state. Error codes hide this explicit transactional behaviour. Notwithstanding, Go as a more recent programming language³ returns an explicit error types instead of raising an exception. This was designed to “encourage you to explic-

³designed in 2007 as a response to problems with C++, Java and Python[1]

itly check for errors where they occur”, instead of “throwing exceptions and sometimes catching them”[5]. Exceptions should contain enough context to find the cause and the source of the error. Exception classes accomplish the grouping of cause source of the error.

In case of error handling with a third-party library, the library called should be wrapped and multiple, possible exception types should be grouped into a unified library exception type. The wrapping reduces the coupling to the third-party dependency. If the library changes (e.g. changing the signature and adding another exception type), only one location has to be adapted to comply to the changed library interface. Additionally, changing to another library is simple and tests can mock the third-party library. For exception handling, wrapping unifies the exception from the library that is clearly distinct from own exceptions.

Another form of exception handling is returning null. Instead of checking for an error code, the caller method checks for a null return value. As easy it is to miss an error code, as easy it is to miss a null checking. Consequently, the program will terminate with an unhandled `NullPointerException` at runtime. Tony Hoare introduced the null reference in 1965 and later called it his “billion-dollar mistake”[6], since it led to many bugs and security vulnerabilities throughout the decades. Languages like Kotlin are designed with Null Safety enforced. Kotlin distinguishes between nullable and non-nullable references, with a compiler enforcing null checking[2]. If a language does not enforce null safety with a compiler, a special case like an empty collection or an optional type can be returned. Passing null to a function as a parameter is also dangerous, since a function would explicitly assert for non-null parameter values. Since this is not a clean way, null parameters should not be used.

2.2.8 Code Smells

Aforementioned guidelines can be rephrased as code smells. Code smells are a characteristic of code that possibly indicate a problem with the clean code rules and therefore maintainability of the system.

2.2.9 Additional Guidelines

Robert C. Martin describes more guidelines for system design, multi-threading, testing and third-party code[7]. These topics are not part of this chapter, since the focus of this work are the aforementioned guidelines and principles.

2.3 Quantitative Metrics for Code Quality

Quantitative metrics express code quality as a quantitative unit. This has multiple advantages:

- A metric sums up the quality of a project in a single unit.
- A quantitative approach tracks the changes in code quality over time. Therefore, it is obvious if code quality improves or not. In case the quality undercuts a threshold, special measures like mandatory refactoring can be undertaken.
- A developers performance can be evaluated based on the code quality. Since the maintainability and reliability of the software depends on the code quality, this is a good incentive to enforce high quality work.
- A certain level of code quality can be required by a contract. As a result, a customer can expect less bugs and a smaller maintenance effort.

Metrics for code quality are a subset of more general software metrics. The following sections describe software metrics that express code quality.

2.3.1 Cyclomatic Complexity

Cyclomatic Complexity is a metric for the complexity of a section of code. It was introduced by Thomas McCabe, Sr. Is is measured by counting the number of linearly independent execution paths. (TODO source)

To compute the Cyclomatic Complexity, the control-flow graph is required. The control-flow graph represents possible execution paths. A node represents a basic block, a code sequence without branching. A directed edge represent jumps in the control flow; e.g. an if-statement without an else-case has two edges. One edge to the basic block in the if-branch and one edge to the basic block after the if-statement. Beside cyclomatic complexity, the control-flow graph is used for compiler optimisation techniques like dead code elimination and loop optimization (<https://ag-kastens.cs.uni-paderborn.de/lehre/material/compil/folien/c101-217a.pdf>).

The cyclomatic complexity on a control-flow graph is defined as:

$$M = E - N + 2P$$

E is the number of edges, N the number of nodes and P the number of connected components. A connected graph is a subgraph, in which all nodes are connected to each other by a path. P represents the number of subprograms (like multiple functions or classes will). Following this definition, the cyclomatic complexity can be calculated.. See figure TODO for a visualisation.

MacCabe recommends to limit the cyclomatic complexity to 10. NIST confirms this recommendation in TODO. A lower cyclomatic complexity improves testability, since the complexity represents the number of execution paths that need to be tested. Therefore, M is the upper bound for number of test cases for full branch coverage. Furthermore, studies suggest a positive correlation between cyclomatic complexity and defects in functions. TODO source Safety standards like ISO 26262 (for electronics in automobiles) or IEC 62304 (for medical devices) mandates a low cyclomatic complexity TODO source.

2.3.2 Halstead complexity measures

Maurice Halstead introduced the Halstead complexity measures in 1977. It is a collection of multiple countable measure and relations.

The Halstead metrics operates on a sequence of tokens that are classified to be an operator or an operand.

The classification enables to count the following properties:

- η_1 as the number of distinct operators
- η_2 as the number of distinct operands
- N_1 is the sum of all operators
- N_2 is the sum of all operands

The total program vocabulary can be computed as:

$$\eta = \eta_1 + \eta_2$$

The program length is not calculated by the lines of code, but as a sum of all operators and operands:

$$N = N_1 + N_2$$

The voume of a program in terms of program length and program vocabulary is defined as:

$$V = N * \log_2 \eta$$

An important metric is the difficulty of understanding a program, defined as:

$$D = \frac{\eta_1}{2} * \frac{N_2}{\eta_2}$$

The major contribution to difficulty is the number of distinct operators η_1 . The combination of difficulty and volumes is the effort for understanding or changing code:

$$E = D * V$$

The effort translates into real coding time following the relation

$$T = \frac{E}{18} s.$$

The number of bugs correlates with the effort following

$$B = \frac{E^{\frac{2}{3}}}{3000}$$

(TODO general critique etc. on this metric since it seems to be arbitrary)

2.3.3 Software Maintainability Index

The software maintainability index was developed by Dan Coleman and Paul Oman in 1994. 16 HP engineers evaluated 16 software systems and scored it in a range from 0 to 100, with 100 representing best maintainability. Following a regression analysis, they identified the following equation to match the maintainability of the evaluated systems (TODO source from Using metrics to evaluate software system maintainability):

$$MI = 171 - 5.2 * \ln \bar{V} - 0.23 * \bar{M} - 16.2 * \ln \overline{LOC} + 50 * \sin \sqrt{2.4 * C}$$

where \bar{V} is the average Halstead Volume, \bar{M} the average cyclic complexity, \overline{LOC} the lines of code and C as fraction of comments.

The software maintainability index was defined many years ago with a limited sample size of developers and projects. Additionally, programming languages have changed significantly over time. A study by Sjøberg et. al suggest no correlation between the software maintainability index and actual maintainability effort in a controlled environment[11]. Although the study only analysis four software projects and lacks generalisation, they find a strong anti-correlation with different maintainability metrics.

Only the code size seems to correlate with the actual maintainability. The former seems to be consistent with a systematic review on software maintainability predictions and metrics by Riaz et. al[9].

Kritik an Clean Code

??Teaching clean code, the paper from one german university

2.4 Tools for Code Quality Analysis

It is good practice to use static code analysis tools to improve code quality. A static code analysis examines a program by analysing an abstract syntax tree, the control or data flow, a pointer analysis or an abstract (approximated) execution. In comparison to dynamic analysis, the code is not executed. The result of a static analysis can be based on approximations, so there might be false positive results. An expert has to examine the result and correct results if necessary[8].

Besides the use of static code analysis for compiler and optimisations, it is also helpful for analysing the code quality. Different categories of code quality related principles can be analysed:

Code Guidelines: A static code analysis can ensure compliance to structural, naming and formatting code guidelines.

Standard Compliance: A requirement for a software may be compliance to an industry standard like IEC 61508 or ISO 26262 (TODO source). Static code analysis can ensure or at least assist with the compliance.

Code Smells: Some code smells follow a known pattern and a static code analysis can detect those smells. Since code smells can be an indication for chaotic code, it is best if these smells are detected and removed early on.

Bug Detection: Although bug detection with static analysis can not detect all bugs, every detected bug before a software release is important. Examples for detected bugs are unhandled expressions or concurrency issues [4].

Security Vulnerability Detection: Static code analysis tools can detect some security vulnerabilities like SQL Injection Flaws, buffer overflows and missing input validation with high confidence. Since these are some common, easy to exploit vulnerabilities, a analysis for security vulnerabilities can increase the security of the overall software [12].

The following sections provide an overview of a few static code analysis tools with focus on code quality and maintainability. I selected them based on popularity and if the projects are still maintained.

2.4.1 PyLint

PyLint is a code analysis tool for the Python programming language. It is open-source and licensed under the GNU General Public License v2.0 and available for all common platforms. PyLint can be executed as a standalone program or can be integrated into common IDEs like Eclipse. A continuous integration pipeline may include PyLint as well to ensure an analysis on every build.

With PyLint, the developer can make sure the code complies to the PEP8 (TODO) style guide for python coding. This includes name formatting, line length and more. It does not calculate a metric for the code but instead warns about violated principles. Additionally, PyLint can detect common errors like missing import statements that may cause the program to crash at start or later at runtime. To support refactoring, PyLint can detect duplicated code and will suggest to refactor the code.

PyLint can be configured to ignore some checks and to disable specific rules. To expand the rule set, a developer can write a "checker", an algorithm to check for a specific rule. The algorithm can analyse the raw file content, a token stream of the source code or an AST representation. The checker can raise a rule violation by providing the location information and the problem type to the PyLint framework and it will be included in the PyLint analysis report.

2.4.2 PMD Source Code Analyzer Project

PMD is an "extensible cross-language static code analyzer" for Java, JavaScript and more. As an open-source project, it is licensed under a BSD-style license and is available for MacOS, Linux and Windows based systems. It integrates into build systems like Maven and Gradle as well as into common IDEs like Eclipse and IntelliJ. PMD can run as part of a continuous integration pipeline and is included in the automated code review tool Codacy (see TODO).

Depending on the target language, PMD supports different rules. For Java, PMD has rule in the following categories:

Best practices: Best practices include rules like one declaration per line or using a logger instead of a *System.out.print()* in production systems.

Coding Style and Design Several rules to improve readability of the code like naming conventions, ordering of declarations and design problems like a violation of the law of Demeter, Cyclomatic complexity calculation and detection of god classes.

Multithreading: Rules to mainly prevent the use of not thread safe objects or methods. Due to the nature of multi-threading and unpredictable scheduling of threads, problems like unpredictable values of variables or dead locks may occur. Since they may not occur in every execution, they are hard to spot and to debug. Consequently, warnings of using non thread safe code may save hours of debugging.

Performance: These rules flag known operations that have hidden performance implications. For example, string concatenation with the plus operator in Java causes the Java Virtual Machine to create an internal String buffer. This can slow down the program execution if numerous String concatenations are performed.

Security: Security wise, PMD only checks for hard coded values for cryptographic operations like keys and initialization vectors.

Error Prone Code: PMD checks for several known code structures that will or may cause a bug at execution. Some rules are part of the clean code principles described in 2.2 and some rules are specific to the target language.

A user can expand the PMD ruleset in two ways: A XPath expression can be specified and will be validated against the AST by the PMD framework. For more control, it is possible to write a Java-based plugin that implements a custom AST traverser. The latter allows for more sophisticated rules and checks.

2.4.3 Codacy

Codacy is a software to "automate your code quality" (TODO source) for more than 30 programming languages. It is available as a cloud-based subscription service with a free tier for open-source projects and a self-host option for enterprise customers. Codacy runs as a cloud software and a user can connect a GitHub, GitLab or Bitbucket repository that will be scanned automatically on every push or trigger a scan of local files with a command-line program. Additionally, a badge can be added to the readme page to show off the analysis results.

Codacy can be seen as a platform that runs multiple different "analysis engines". These analysis engines combine multiple tools depending on the language. For Python,

Codacy uses PyLint, Bandit (a scanner for security issues) and a metric plugin. Due to the licensing of those analysis engines, the engines are open-sourced, whereas the Codacy platform is proprietary.

Depending on the language, Codacy supports scanning for common security and maintainability issues. The latter is faced with scanning for Code Standardization, Test Coverage and Code Smells to reduce technical debt.

Codacy allows customization by disabling certain rules and changing rule parameters like the pattern to fit the rules to the project's preference.

2.4.4 Sonarqube

Sonarqube is an analysis tool to maintain high code quality and security. It supports 15 programming languages in the open-source version licensed under the GNU Lesser General Public License, Version 3.0. In the Developer Edition, Sonarqube supports 22 languages and 27 languages in the enterprise edition. Sonarqube is a self-hostable server application with a SonarScanner client module that can be integrated in build pipelines like Gradle and Maven as well as a command-line tool for other build pipelines. The SonarScanner reports the result to the server, that serves a website to review the results. With the Developer Edition, branch and pull request analysis are possible and a pull request will be annotated with the analysis results.

For Python, Sonarqube offers more than 170 rules. These rules are in the following categories:

Code Smells: Code Smells like duplicated String literals are flagged with four different levels of severity (from higher to lower severity): Blocker, Critical, Major, Minor. Explanations and examples are accessible for the developer to understand and fix the issue. The analysis is more in-depth than AST-based solutions, since it additionally uses control and data-flow analysis.

Bug: Multiple rules cover bugs that would definitively result in a runtime exception and program termination. As an example, calling a function with the wrong number of arguments will be flagged with the highest severity, since it will raise a `TypeError` during runtime. Although programmers may notice issues like this during coding and testing, the issue can remain hidden if it only triggers in a special execution path of the system.

Security Hotspot: The Security Hotspot analysis unfolds pieces of code that may be

a real vulnerability and requires a human review. The scanning has a hotspot detection for seven out of the OWASP Top Ten Web Application security risks (TODO src). A flagged hotspot contain a detailed explanation of the reason for being flagged and a guide on how to review this hotspot. The recommendation to double check a piece of code can help to prevent a security vulnerabiliy from being deployed to production.

Security Vulnerabilities: A security vulnerabilites analysis reveal code that is at risk of being exploited and has to be fixed immediately. Especially misconfiguration of cryptographic libraries can be revealed easily and future damage prevented.

Additionally, Sonarqube offers several metrics like test coverage and an custom derivate of Cyclomatic Complexity named Cognitive Complexity. The authors sees several shortcommings in the original Cyclomatic Complexity model:

- Parts of code with the same Cyclomatic Complexity do not neccessarily represent equal difficulty for maintenance.
- Cyclomatic Complexity does not include modern language features like try/catch and lambdas. Therefore, the score can not be accurate.
- Cyclomatic Complexity lacks useability on a class or module level. Since the minimum complexity of a method is one and a high aggregated Cyclomatic Complexity for a class is measured for small classes with complex methods or large classes with low complexity per method (e.g. a domain class with just getter/setter).

Cognitive Complexity adreses these points, especially the incorporation of modern language features and meaningfulness on class and module level.

The calculation is based on three rules:

Ignore Shorthand Structures : Method calls condense multiple statements into one easy to understand statement. Therefore, method calls as shorthand are ignored for the score. Similiarly, shorthand strcutures like the null-coalescing operator reduces the Cognitive Complexity compreaed to an extensive null check and is therefore ignored for score calculation.

Break in the linear control flow : A break in the expected linear control flow by loop structures and conditionals adds to Cognitive Complexity as to Cyclomatic Complexity. Additionally, Catches, Switches, sequences of logical operators, recursion

and jumps also adds to Cognitive Complexity, since these structures break the linear control flow.

Nested flow-breaking structures : Flow-breaking structures that are nested additionally increase the Cognitive Complexity, since it is harder to understand than non-nested structures.

The method-level Cognitive Complexity score represents a relative complexity difference between methods that would have the same Cyclomatic Complexity. Furthermore, the aggregated Cognitive Complexity on a class or module level now differentiate between classes with many, simple methods and classes with a few, but complex methods. As a result, domain classes (containing mainly getters/setters) have a lower score compared to classes with complex code behaviour. (TODO source whitepaper multiple times)

Developers can extend Sonarqube similarly to PMD. They can write a Java plugin that can access the AST but also a semantic model of the code. Additionally, extensions can be created that provide functionality to other extensions. The Java plugin is compiled into a .jar file and placed into the plugin dir of the Sonarqube installation. For simpler rules, XPath expressions are possible through the web interface and allow a quick extension. For instance, if developers observe bad code style during a pull request review, they can quickly write a rule to enforce this rule in all subsequent pull requests automatically.

- Integration in continuous integration, manual review

- Tools review for all tools that check different stuff

- Single Responsible principle, Open-Closed principle

3 Approach

Approach(es) without the quantitative results (or only selected results to motivate the choices/the problems)

- What has possibly failed (short)

- What has worked and why

- What could be possible method extensions/improvements

This section is divided into the approach for the Clean Code Analysis Plattform and the Clean Code classification.

3.1 Clean Code Analysis Plattform

The goal of the design and implementation of the Clean Code Analysis Plattform (CCAP) is a tool for software developers to improve the code quality of existing and new code. The tool accepts a directory containing source code files as input and analyses the input for snippets of improvable code quality. If the analysis classifies a code snippet as problematic, it should help the developer to improve the snippet by providing information about the problem. Ultimately, this should train the developer to spot problematic code by its own and to write clean code by default, so the number of alerts should decrease. At the same time, the overall software quality of a project increases immediately at rewriting a marked snippet and in the long term at training the developer to write code with higher quality.

In order to use the tool effectively, the design and implementation should cover the following requirements:

Useability : The CCAP should be an easy-to-use tool. Developers shall be able to install and run the tool. The extra effort of using this tool should be small, and the developer should use the tool in his day-to-day workflow without additional friction. The developers can interpret the issue and localise the problematic code spot immediately.

Expandability : The extension of the detected code problems should be easy. A clear defined interface for extensions is required so an extension developer would not need specific knowledge about the internal architecture of the tool. The expandability allows the desired workflow of a developer finding problematic code in an, e.g. peer-review, formalising it into an extension and sharing this extension with the team. With each iteration, the code quality of all team members would increase.

Integration : The tool should be easy to integrate into different systems. These systems include local workflows like git pre-commits or build systems and remote continuous integration/delivery/deployment pipelines.

A more specific requirement is Python as an input language and the expansion language. After JavaScript, Python is the second most popular programming language 2019 according to the Github statistics (<https://octoverse.github.com/#top-languages>). Besides the general popularity, Python is heavily used in the scientific community for machine learning and universities for teaching programming. These groups are part of the potential target audience, and students in specific can benefit from automated reporting of low-quality source code.

3.1.1 Architecture

The CCAP architecture is divided into a static part and two extension possibilities: An extension with analysis plugins adds more rules that are validated by the system. Adding an output plugin allows specifying the output format to fit custom workflow needs. The static part consists of four components: A core component to act as an orchestration unit, a component for handling the source code input and a component for handling the analysis plugins as well as output plugins. The design follows the requirements and goals for the platform. (TODO: High-level architecture schematic diagram)

The core component contains the main function and handles argument validation and parsing. Furthermore, it orchestrates other components by initialising and executing those. This process is divided into the argument parsing, initialisation and execution phase: In the argument parsing phase, the command line arguments are parsed and validated. It validates the existence of the required input directory argument and the optional plugin path configuration for the analysis plugins and the output plugins. Additionally, the logging level and the output format can be defined. The latter determines, which output plugin will be used, although the existence of the specified plugin is not

validated in this phase. A parsing or validation error will cause a program termination without further processing. The initialisation phase instantiates all components and the analysis plugin handler will scan the specified directories for plugins and keeps an index of all found plugins in memory. The output plugin manager scans for an output plugin, that satisfies the specified argument. If no plugin matches the output argument, the program will terminate with a one exit code and a failure message to indicate the problem. In the execution phase, the input handling component scans the input directory for files ending with *.py* and parses the source code into an Abstract Syntax Tree (AST) per file. In the next step, the core passes the parsed data to the analysis plugin handler. The latter will execute all plugins for all files and collects the results. Afterwards, the core component calls the output plugin component to output the results. If no exceptions occur during the execution phase, the program will be terminated with a zero exit code indicating the successful run.

The input component scans the given input directory for all Python source code files and parse the source code into an AST. For scanning the input directory, an algorithm will walk recursively over all folders and files. The detection of Python source code files is based on the file ending *.py*. The algorithm will return a list of file paths and the corresponding file content. Next, the AST parser is called and will add a parsed AST object to the list besides the file path and content. This list will be passed to the analysis plugins by the analysis plugin handler. An alternative approach would be to not read and parse the code in the input component, but instead, let the plugins read and parse the file content if needed. With many files to scan, the latter approach would have a lower memory footprint since the file content and the AST will not be held in memory. Consequently, every analysis plugin has to perform an expensive read operation from disk and the performance scales with the number of files and the number of analysis plugins. If the input component reads all files, the information is held in the main memory and the performance only scales with the number of files. Since the files are text-based, the number of files needed to exceed the main memory is expected to be high enough to fit most projects. (TODO sample calculation?).

The analysis plugin handler manages all analysis plugins. This component finds all plugins, executes the plugin and collects the reported results. During the initialisation phase of the core component, the analysis plugin handler will scan the plugin directory for all Python files. It imports all Python files and scans those for classes, that inherit from an abstract *AbstractAnalysisPlugin* class. The abstract class defines all methods that need to be implemented in the concrete plugin subclass. The analysis plugin handler

instantiates all found classes. During the core execution phase, the core receives a list with the file name, file content and the parsed AST. All plugins are called on a specific entry point method that is defined in *AbstractAnalysisPlugin* with the aforementioned list. The plugin will return an instance of *AnalysisReport* with the plugins metadata and a collection of problems. The report is collected for every plugin into an *FullReport*. Additionally, the report contains information about the overall plugin execution time and run arguments. After all plugins have been executed for all files, the analysis plugin handler returns the *FullReport* to the core component.

The last component in the execution chain is the output plugin handler that will pass the *FullReport* to the specified output plugin. It implements the same algorithm as the analysis plugin handler to find all plugins that inherit from *AbstractOutputPlugin*. Instead of keeping track of all plugins, only the plugin that corresponds to the output format argument is instantiated. The output plugin has an entry point as defined in *AbstractOutputPlugin* that is called with all the collected results.

Analysis Plugins

Analysis plugins provide the easy extendability of the platform by developers. All users of the tool can expand the set of problems it can detect by implementing a plugin Python and placing it into the plugin directory of the tool. In order to be compatible with the core components, a plugin has to be a class that inherits from the *AbstractAnalysisPlugin* class. Firstly, the abstract class introduces a class member variable *metadata* of the class *PluginMetaData*. A concrete plugin class sets this class member in the constructor to provide a plugin name, author and optional website. This metadata is used in the output to show more information about the plugin that reported a problem. Secondly, *AbstractAnalysisPlugin* specifies a *do_analysis* method that serves as a entry point that the platform will call. It accepts a list of *ParsedSourceFile* objects that contains the file path, the file content and the corresponding AST for all input files. With this information, the plugin can implement any logic to detect problems. For example, the plugin can traverse the AST to look for specific node types, it can use a tokeniser on the content of the files and analyse the token stream or it can run sophisticated machine learning algorithms on the source code. The plugin can even import third-party libraries, although they have to be installed by the user on the system. After detecting all problems, the plugin returns an *AnalysisReport*. The report contains the plugins metadata and a list of found problems. These problems are instances of a problem class inheriting from *AbstractAnalysisProblem*. The abstract class expects a

file path and line number as constructor arguments and requires the plugin developer to override the problem name and description (see listing ?? for an example). The problem name and description will be shown in the final output and should follow the following guidelines:

```
class ReturnNoneProblem(AbstractAnalysisProblem):
    def __init__(self, file_path, line_number):
        self.name = "Returned None"
        self.description = "Returning None is dangerous since the caller has to check for None.
Otherwise, a runtime exception may occur."
        super().__init__(file_path, line_number)
```

Listing 3.1: Example for overwriting the *AbstractAnalysisProblem* with a concrete implementation. The problem name and description are overwritten.

- Have a proper name that allows the experienced developer to recognise the problem quickly.
- Explain what code construct is problematic.
- Give reasons why this code is seen as problematic.
- Show guidance and examples on how to fix the problem and improve the code.

Although it is possible to use one plugin for multiple, different problem types, having one plugin for one problem type helps to reuse and share the plugin. Additionally, it can be disabled easily by removing the plugin from the plugin folder.

Steps to create an analysis plugin In order to expand the CCAP with an analysis plugin, the following steps are required for a developer:

1. Create a *.py* file with a class inheriting from *AbstractAnalysisPlugin*.
2. Instantiate *PluginMetadata* and assign it to the metadata member.
3. Define a problem class inheriting from *AbstractAnalysisProblem* and set the problem name and description following the guidelines above.
4. Implement the *do_analysis* method with a *ParsedSourceFile* parameter. Return an *AnalysisReport* instance with all found problems.
5. Place the *.py* file into the analysis plugin directory of the tool.

Whereas these are the minimum required steps to implement the plugin, the developer is free to add additional methods, classes or import libraries as necessary. Furthermore, it is advisable to implement several tests to ensure the plugins correctness. CCAP uses `pytest`¹ for testing, implementing a test is as easy as writing a function with a “test_” prefix and running the `pytest` command. The plugin can be tested by simulating a call from the CCAP core with a mocked `ParsedSourceFile`.

Return None Plugin A rather simple analysis plugin demonstrates the capabilities of the CCAP: The Return None Plugin scans the source code for functions that return the `None` value. Returning `None` may result in runtime exceptions if the function caller does not expect and handle a potential `None` return. Although the `None` can be returned directly or as a value of the returning variable, this plugin only focus on the explicit `return None` statement to showcase the options for the developer to analyse the source code.

Detecting `return None` is possible in multiple ways. The following shows the possibilities for developer to implement such a detector:

Regular Expression : In the `do_analysis` method, a developer has access to the source code as a string. Therefore, it is straitforward to use a regular expression to detect a `return None` statement.

```
import re
matches = re.finditer(r"return None", source_file.content, re.MULTILINE | re.DOTALL)
```

Since the regex library only returns the start and end index of matches, these have to be converted to line numbers. Afterwards a `ReturnNullProblem` instance can be created for every match and added to the `AnalysisReport` for this plugin.

This approach uses regular expressions to match patterns in a string without utilising the structures of the code. It is a simple but powerful way and most developers are familiar with regular expression. On the flip-side, source code may have various ways syntactic ways to express a semantic. Consequently, the regular expressions have to be designed carefully to cover all variations. For instance, it is possible to encounter a doubled whitespace, that would break the aforementioned regular expression. Additionally, regular expressions do not operate on the structure of the code; therefore, they can not detect high-level patterns on the code structure.

¹<https://docs.pytest.org/en/stable/>

Tokenisation : The process of dividing a character stream into a sequence of tokens is called tokenisation (also known as lexical analysis). With the Python tokeniser, a token can contain multiple characters and has a token type like name, operator or indentation. A token sequence provides more information about the code structure that can be used to detect problematic patterns. A token sequence for a *return None* statement would be the following:

```
[...( type: NAME, value: return), (type: NAME, value: None)...]
```

A simple algorithm would scan the sequence for two subsequent name tokens with the values "return" and "None". The abstraction level of a token sequence is higher than of a character sequence. Since the whitespace between "return" and "None" provide no semantic meaning, it is removed on this abstraction level. Whereas the regular expression based approach would have to deal with problems as the doubled whitespace, the token-based approach profits from the higher abstraction level and can access meaningful tokens directly.

Abstract Syntax Tree : After tokenisation, a parser takes the token stream and prases it into a hierarchical data structure like an Abstract Syntax Tree. With an AST the code is represented structurally and it is possible to traverse the tree following the structure of the code. The AST consists of nodes with different type and children, for example, a node representing an if statement. It is possible to access the condition and the body of the if statement as descendant nodes in the AST. Analysing an AST allows a higher abstraction level then analysing the token string since the code structure is represented. Therefore, more abstract rules that analyse the code structure are possible.

To detect a *return None* statement in the abstract syntax tree, an algorithm would traverse all AST nodes looking for a return-typed node. If found, the value descendant contains the statement part after the return keyword. A *None* value would be represented as a node of type constant or name constant (depending on the parser version). The value descendant of the constant node may then be checked for equality to *None*. All AST nodes contain the corresponding line number, that are necessary for creating a problem message. See listing 3.2 for a Python implementation.

Analysing the AST is best for problems in the code structure since the AST represents the code structure in a well-defined, traversable data structure. Although

simpler patterns like the *return None* could be detected using regular expressions, a detection on AST level allows detecting the semantic meaning instead of the syntactic representation of a problem. For more complex problems, it is inevitable to use the AST most of the time.

```
for node in ast.walk(a.ast):
    if isinstance(node, ast.Return):
        return_value = node.value
        if isinstance(return_value, ast.Constant) or isinstance(return_value, ast.
NameConstant):
            if return_value.value is None:
                problems.append(ReturnNullProblem(a.file_path, return_value.lineno))
```

Listing 3.2: Detecting RETURN NONE problems by analysing the AST data structure.

This implementation covers only the simple case, in which *return None* is written explicitly. Of course, several modifications are possible, that would not be detected by this plugin and would require more sophisticated algorithms. Evaluating, if machine learning models trained on this simple rule can also detect such modifications is part of the second part in research question 3 (TODO: check, [jump link](#)).

Since detection alone is not a great help for the user, a useful description of the problem is necessary. As described in section 2.2.7, there are some alternatives to returning none. If none happens due to an error, it is better to raise an exception and provoke an explicit error handling, which prevents runtime errors. If the standard return type is a collection, an empty list is more appropriate, since the function caller will most likely write the logic for an unknown amount of list items. However, if the standard return type is a single object, it is not apparent what to return. With PEP484 (TODO source), an optional type was introduced for type hints in Python 3. Using a static type checker for Python like mypy (TODO source) outputs a warning if a developer forgets to check an optional for not being none.

Condition with Comparison The second plugin detects a direct comparison in a conditional statement. For better readability and understandability, it is suggested to use a function call with an appropriate name that evaluates to true or false. This function call allows a natural reading of the conditional statement without deciphering the meaning of the boolean logic.

An if-statement consists of a condition and body part. If the condition evaluates to true, the body is executed. For the condition, any logical expression will be evaluated. A simple algorithm would take the AST, search for if nodes and check if the condition part

is a compare node. However, negation and logic AND/OR would not be detected by the simple approach. Consequently, a recursive algorithm has to follow logic operators and checks the expressions for comparison nodes. The simple approach does not provide value for the CCAP but we use it for the machine learning evaluation with the problem type name `CONDITION COMPARISON SIMPLE`.

Following these considerations, an advanced algorithm would scan for if nodes in the AST. The conditional part as a boolean expression is evaluated in a recursive function that returns true if a direct comparison is made anywhere in a boolean expression. The conditional part can be a boolean operation (like AND/OR) or a unary operator (like NOT). In these cases, the recursive function will be called again with the respective expressions. If the conditional part is neither a boolean operation nor a unary operator, a true is returned if the conditional part is not a method call. The last case would be the base case in the recursion. Listing 3.3 shows the Python implementation of the recursion.

```
def _check_if_direct_comparison(self, node):
    if isinstance(node, ast.BoolOp):
        violated = False
        #check all expressions of the boolean operator
        for value in node.values:
            if self._check_if_direct_comparison(value):
                violated = True
        return violated
    elif isinstance(node, ast.UnaryOp):
        return self._check_if_direct_comparison(node.operand)

    return not isinstance(node, ast.Call)
```

Listing 3.3: Recursive function to analyse an if statement for direct comparisons. Since a condition should contain a method call, the function returns False if this is not the case.

Output Plugins

With output plugins, CCAP adds additional flexibility towards the output format. Depending on environmental requirements, the output can be adapted with custom logic by defining an output plugin. For instance, running the tool locally could write the results to the standard output in a human-readable way or create a formatted HTML file to be displayed in the browser. A machine-readable JSON output may be preferred

when running inside an automated workflow.

Output plugins follow similar concepts as analysis plugins. A plugin class inherits from *AbstractOutputPlugin*. The abstract class introduces the metadata member of the *PluginMetaData* class to encapsulate plugin name and author information. Additionally, a second member variable *output_format* has to be defined with a short abbreviation for the output format. This field is used by the output plugin handler to select the output plugin by the run arguments. Consequently, it should be unique; otherwise, the first found plugin with a matching *output_format* field is selected by the output plugin handler. Therefore, using a custom prefix string is recommended.

As entry point method, the method *handle_report* has to be implemented. The method provides an instance of *FullReport* as its argument. The *report* field holds a collection of *AnalysisReport* for every analysis plugin that has been executed. With metadata and problem information, the output plugin has access to all required information to produce the desired output.

Steps to create an output plugin To expand the output capabilities of the CCAP, the following steps are, similar to analysis plugins, required:

1. Create a *.py* file with a class inheriting from *AbstractOutputPlugin*.
2. Instantiate *PluginMetaData* and assign it to the metadata member.
3. Set the *output_format* field with a unique abbreviation for this output format. Since it should be unique, a custom prefix prevents a name collision with pre-existing plugins.
4. Implement the *handle_report* method with a *ParsedSourceFile* parameter.
5. Place the *.py* file into the output plugin directory of the tool.

Standard Output Plugin The Standard Output Plugin writes the formatted output to the stdout stream. It is enabled if no output plugin is explicitly specified.

The output is divided into a general part and the problem list. The former contains the input path, all executed plugins, a total runtime and a summary field with the numbers of problems in total. The latter displays the list with problems, grouped by analysis plugin. For each problem, the problem name, file path, line number and description is printed. A colon separates the file path and line number. Some terminals parse the path

and line number so that a user can click the path and it opens the default editor with the cursor in the correct line. Sample output is shown in listing 3.1.1.

```
Analysis Report on /Users/d064518/MA_CleanCodeAnalyser/test_programs.
Analyse Plugins: Condition Method Call Plugin, Simple Condition Method Call Plugin, Return None
  (Null) Plugin, Sample Plugin.
Total time: None
Summary: Found 16 problem(s)
-----
PROBLEMS:
  PLUGIN NAME: Condition Method Call Plugin by Enrico Kaack <e.kaack@live.de>
    Found: Explicit comparison in condition in /Users/d064518/MA_CleanCodeAnalyser/
test_programs/return_none.py:18
      Explicit comparisons in conditions should be replaced by method call for better readability
...
  PLUGIN NAME: Return None (Null) Plugin by Enrico Kaack <e.kaack@live.de>
    Found: Returned None in /Users/d064518/MA_CleanCodeAnalyser/test_programs/
return_none.py:4
      Returning None is dangerous since the caller has to check for None. Otherwise, a runtime
exception may occur.
```

HTML Output Plugin For larger projects, the Standard Output Plugin may not be sufficient to get an overview of the project or to generate a report. Therefore, the HTML Output Plugin creates a *report.html* file, that can be opened in every browser and provides a cleaner overview. It has similar general data at the top as the Standard Output Plugin, but it moves the problem description into a tooltip to offer a more compact overview. See figure ?? for an example screenshot.

3.1.2 Extensions

While the core functionality exists, extensions for the CCAP would aim for improved useability for the user. One important improvement in useability would be an IDE integration for common IDEs like Visual Studio Code. A good starting point would be the Language Server Protocol². This IDE integration requires a language client as a Visual Studio Code extension and a protocol-compliant language server. The latter would wrap the CCAP and modify it slightly to process a single, changed document and return the results. Since the CCAP architecture is modular, this is possible without larger modifications.

²<https://microsoft.github.io/language-server-protocol/>

Analysis Report on /Users/enrico/MA_CleanCodeAnalyser/test_programs	
Loaded Analysis Plugins	Condition Method Call Plugin, Simple Condition Method Call Plugin, Return None (Null) Plugin, Sample Plugin
Summary	Total Time: None 16 problem(s) found
RESULTS	
Condition Method Call Plugin by Enrico Kaack	
Explicit comparison in condition	Explicit comparisons in condition statements should be replaced by method call for better readability. in /Users/enrico/MA_CleanCodeAnalyser/test_programs/return_none.py:18
Explicit comparison in condition	Explicit comparisons in condition statements should be replaced by method call for better readability. in /Users/enrico/MA_CleanCodeAnalyser/test_programs/return_none.py:3
Explicit comparison in condition	Explicit comparisons in condition statements should be replaced by method call for better readability. in /Users/enrico/MA_CleanCodeAnalyser/test_programs/return_none.py:12
Explicit comparison in condition	Explicit comparisons in condition statements should be replaced by method call for better readability. in /Users/enrico/MA_CleanCodeAnalyser/test_programs/condition_without_method_call.py:40
Explicit comparison in condition	Explicit comparisons in condition statements should be replaced by method call for better readability. in /Users/enrico/MA_CleanCodeAnalyser/test_programs/condition_without_method_call.py:40

Figure 3.1: Output of the HTML Output Plugin, displayed in a browser.

Another feature would be a configuration to disable specific analysis plugins, that one does not want to run on the project. With the current architecture, this would be possible by moving the analysis plugin file outside the plugin directory. A configuration with command-line parameters or with a per-project, hidden config file would have a better user experience. The latter could also be committed to the version control system, so every team member has the same configuration. Continuing on configuration possibilities, having configuration possibilities for plugins would increase flexibility for a plugin developer. Introducing configurable warn levels like ‘warning’ or ‘error’ could help in continuous integration pipelines to decide if a build succeeds but the warnings are reported or if a build should fail because of error-level problems. Some rules could be seen as recommendations (warning level), whereas other rules would be unacceptable (error level).

Lastly, a feature to disable problem reporting for a specific code location could bring a boost in user acceptance. While some clean code rules are objective and can be measured precisely, some rules are more general recommendations that may not apply to every occurrence of this situation. For instance, the clean code guidelines generally suggest not to have more than three function arguments; it may be necessary or even inevitable to have four arguments. The CCAP should report this as a problem, but

the user should be able to decide if it is acceptable. In this case, the problem type on this specific location should not be reported in the future. In the current version, the user has no way to ignore a specific problem. Consequently, over time the number of problems the user has to ignore will accumulate until the user abandons the tool since it does not provide an added value over the frustration of manually ignoring problems.

3.2 Clean Code Classification

In the previous chapter, we have presented a platform to check for clean code rules. This platform requires analysis plugins to detect certain problems. These handwritten rules work well on certain rules, but certain rules may require a complicated analysis of the AST, the data-flow or the inter-module dependencies. Some rules on the other hand are subjective and we do not see a way of detection those rule violations with an algorithm. Therefore, this chapter introduces an approach to evaluate different supervised machine learning models to classify code into clean code or problematic code. Furthermore, we describe how we modify the code to test the generalisation capabilities of the models. The mentioned machine learning models include random forest classifier, support-vector-Machines, gradient boosting classifier and a recurrent neural network based on Long short-term memory units.

The objectives are to train, evaluate and compare Source Code Snippet Classification models and to modify the input data to represent a problematic code in an unseen way.

In the following we describe the approach in detail. The structure follows a machine-learning typical pipeline of data collection, preprocessing, encoding, splitting and training.

3.2.1 Challenges

Lack of datasets The most crucial components of every machine learning experiment are labeled datasets. For clean code detection, our research has shown no datasets with labeled clean code violations that could be used. As a result, we evaluated different approaches:

Solution 1: Related research fields like code completion often use the py150 (TODO source) dataset to train models. This dataset contains 150.000 python files, sourced from GitHub. For supervised classification tasks, this dataset is inadequate due to its missing labels.



Figure 3.2: Example commit messages from different repositories. Those examples highlight (1) the missing naming convention, (2) inclusion of additional performance improvements and (3) several edits in multiple files. Those commits are found by searching GitHub for “clean code improvements” and filtering for Python language and commits. Sources: ¹, ², ³

Solution 2: Another possible data source would be git commit histories. We could scan the history for commits that fixed unclean code. The previous code could then be labeled “unclean” and the committed code would be “clean”. Same applies to issues and referenced pull requests on GitHub. We dismissed this approach for several reasons (see figure 3.2 for an illustrative example):

1. There is no annotation in neither issues descriptions nor git commit messages that would reliably tell, if the change is changing a violation of the clean code rules.
2. Even if we would be able to find commits that improve chaotic code, we could not ensure automatically, that the commit message is correct and only the improvement is included in the commit.
3. Searching for commits, we found several clean code improvements and refactorings of large portions of the code base in one commit. Additionally, it is often not explicitly mentioned, which rule applied for the improvement. Consequently, a training with such data could only lead to binary classification into clean code or unclean code. From a practical standpoint, it is advantageous to name the explicit rule violation and to explain how to improve.

¹https://github.com/Zadigo/ecommerce_template/commit/db2655c7ba16add17db221761fabda28e0128761

²<https://github.com/gbaranski/homeflow/commit/7d06d9f6ea78530e5a542f71f6ddd536e2d55bb7>

³<https://github.com/ptimofeyeff/distributed-computing/commit/>

Label	RETURN NONE	CONDITION SIMPLE	CONDITION COMPARISON
0	99.79%	97.36%	95.14%
1	0.21%	2.64%	4.86%

Table 3.1: Class distribution for the train/test dataset before the split. Label [0] represents clean code and label [1] marks problematic code of the corresponding category.

Solution 3: Although Python code is available in abundance in form of open-source projects, it is possible to manually label source code that violates clean code guidelines. For the following reasons, we dismissed this approach:

1. Curating a hand-labeled dataset is a lot of work that takes time.
2. Manual labeling does not ensure correct labeling. A human can make mistakes or can subjectively misinterpret chaotic code as clean code. The quality of the dataset could suffer and limit the machine learning models performance.

Solution 4: Given any amount of Python files, we could use the analysis plugins from section TODO analysis plugin to reliably detect violations of the corresponding rule. Based on this labeled data, we can train the classifier to detect these rules. If there would be a labeled dataset for more complex clean code rules, we could transfer the findings to these more complex rules. We choose this solution as to solve to solve the dataset problem and train classifier, with more details given in chapter TODO. As a second step, we manipulated the input data, so the original analysis plugin could not detect the still existing problems anymore. With the second step, we show that the classifier generalise well enough to may work on complex clean code rules.

Inbalanced Dataset ?? A second challenge is a major imbalance between the classes labeled as problematic code and clean code. For the three different problem types, RETURN NONE, CONDITION COMPARISON SIMPLE and CONDITION COMPARISON, the class distribution is shown in table 3.1.

Solution 1: Due to the inbalance, we do not choose accuracy as a suitable evaluation metric, since classifying all samples as clean code would result in a accuracy of 99.79% for the RETURN NONE problem type. Instead, we use precision and recall as metric,

ff5e1322611a0c3b3c62556d7a775441630df252

with a f1 score as single, weighted metric. More details about the metrics are in chapter TODO.

Solution 2: We will apply undersampling and oversampling techniques on the data and evaluate the difference in model performance. Undersampling removes random samples from the majority class and thus balance class labels. Oversampling replicates random samples from the minority class to balance class labels while also increasing the overall amount of data. This technique is only used on the trainings data. Since it changes the data distribution, an additional train-dev dataset is used as an additional test set to evaluate the impact of the resampling on the model performance.

Solution 3: Our approach of training different classifier is a solution to the data imbalance, since different classifier have a different sensitivity to data balance.

SVM quadratic complexity in training Support-Vector Machines have a quadratic time complexity on the number of trainings data. Consequently, the training takes a lot of time and the oversampling strategy worsen the trainingtime additionally.

Solution: We reduced the number of experiments with SVM and just evaluate some simple cases for baseline performance.

source: Time Complexity Analysis of Support Vector Machines (SVM) in LibSVM

3.2.2 Dataset

The source of our dataset are open-source python projects on GitHub. We queried the top starred python repositories and handselected 18 repositories. Although most top starred repositories are data science frameworks, we additionally choose projects from domains such as web server, automation and containerisation. See table TODO for a list of all 18 repositories. A script downloaded all projects in their main branch with the current head. See table TODO attachment for the corresponding git hashes. Afterwards, we removed all non-python files. Next, we uniformly sampled 20% of all files and separated those into a holdout set for final testing. Additionally, we perform our train/test split on the file level and not on the sample level. As described later in chapter 3.2.3, we use a sliding window approach to convert source files with dynamic length into fixed length samples. Due to the sliding window approach, we may have one problematic code line in multiple samples. If we would do a train/test split on sample level, we may

3 Approach

Parameter								RN			CCS			CC		
over-samp.	under-samp.	emb-size	epochs	batch-size	#lstm-cells	dropout		Train	Train/Test	Test	Train	Train/Test	Test	Train	Train/Test	Test
-	-	32	3	256	10.0	0.2 0.2	<i>F1</i>	0.9937	0.9931	0.9927	0.9935	0.9817	0.9825	0.9916	0.9799	0.9812
							<i>Rec1</i>	0.9906	0.9948	0.9912	0.9957	0.9864	0.9875	0.9925	0.9811	0.9832
							<i>Prec</i>	0.9968	0.9913	0.9941	0.9912	0.9769	0.9776	0.9908	0.9787	0.9791
-	0.5	32	3	64	10.0	0.2 0.2	<i>F1</i>	0.9995	0.8107	0.8072	0.9987	0.9699	0.9713	0.9978	0.9688	0.9699
							<i>Rec1</i>	1.0	1.0	1.0	0.9985	0.9968	0.9971	0.9988	0.9968	0.9978
							<i>Prec</i>	0.999	0.6817	0.6768	0.9989	0.9445	0.9469	0.9968	0.9424	0.9436

Table 3.2: TODO lstm

introduce samples covering one problematic line at different positions in both datasets (see figure TODO for visualisation). This would diminish the validity of metrics calculated from the test dataset. Table 3.2 shows an example for a train and test metrics of an neural network with LSTM cells. Without over- or undersampling, we achieve near perfect results for all problem types. If we apply undersampling to the train data, we basically remove some samples labeled as negative (in this case containing clean code). We observe a drop in precision but a near perfect recall. A decrease in precision and a constant recall means an increase in false positives. Since we removed some negative samples with undersampling, the overfitted model has not seen those samples that exist in the test set. Consequently, we decided to perform a train/test split on file level as described earlier.

We ensured to have a similiar data and label distribution on all data sets. Table 3.3 shows general and problem specific metrics on file level for the train, test and holdout set. The problem specific metrics are collected after processing the code files as described in chapter 3.2.3. We define the metrics as follows:

We calculate the *average lines of code per file* for n files with l_i as the lines of code for file i with the following equation:

$$\text{average LOC per File} = \frac{\sum_{i=0}^n l_i}{n}$$

For the *proportion of lines of code containing a problem*, we assume one line contains the problem. Therefore, we define this metric for n files with l_i lines of code per for the i th file and p_i problems in the i th file as follows:

$$\text{LOCs containing problem} = \frac{\sum_{i=0}^n p_i}{\sum_{i=0}^n l_i}$$

Last, we calculate the number of problems per file for n file and p_i problems in the i th

	Metric	Train	Test	Holdout
	Lines of Code	2,530,455	246,813	671,554
	Number of Files	13,330	1,481	3,702
	average LOC per File	189.83	166.65	181.40
Return None	LOCs containing problem	0.07%	0.09%	0.08%
	Problems per File	0.14	0.15	0.14
Condition Comparison Simple	LOCs containing problem	1.29%	1.21%	1.32%
	Problems per File	2.44	2.02	2.4
Condition Comparison	LOCs containing problem	2.34%	2.31%	2.44%
	Problems per File	4.44	3.85	4.42

Table 3.3: General and problem specific metrics for the train, test and holdout set. The lines of code containing a problem and the problems per file are similar.

Problem Type	Label	Train	Test	Holdout
Return None	[0]	99.79%	99.73%	99.78%
	[1]	0.21%	0.27%	0.22%
Condition Comparison Simple	[0]	97.34%	97.49%	97.23%
	[1]	2.66%	2.51%	2.77%
Condition Comparison	[0]	95.13%	95.21%	94.9%
	[1]	4.87%	4.79%	5.1%

Table 3.4: Label frequency for all problem types on the train, test and holdout set.

file as:

$$Problems\ per\ File = \frac{\sum_{i=0}^n p_i}{n}$$

The most important metrics for the data distribution is the proportion of lines of code containing the problem type and the average number of problems per file. For the first, we observe a maximal difference of 0.02 percentage points for the RN problem type, 0.11 percentage points for CCS and 0.13 percentage points for the CC problem type. This directly translates into the label distribution on sample level shown in table 3.4 with a maximal difference 0.06, 0.26 and 0.31 percentage points for the three problem types. The number of problems per file is lower in test set for the CCS and CC problem type. Nevertheless, since the first metric and the label frequency are comparable, we assume our datasets having a similar data distribution.

TODO per project analysis table in attachments.




Figure 3.3: Split of the dataset. First, 20% of all files are separated into a holdout set for later testing. The remaining 80% of the files are split into 90% training and 10% test data.

3.2.3 Processing

For the processing steps, we create a pipeline with the d6tflow framework³. Each processing step is a task, that stores its results depending on the parameter configuration in a pickle file. This allows to define a pipeline once and run it with different parameter combinations. The scheduler automatically determines what tasks to run and what task outputs are already stored in pickle files. Consequently, running the pipeline is more efficient and repeatable.

The processing pipeline consists of tasks to read the files into a datastructure, to process source code and find the problems, to create a vocab dictionary, to convert the data into labeled samples and to split the data into a train, traintest and test set. See overview TODO for a schematic representation of the pipeline.

Files into internal Data Structure

First, a scanner walks recursively in all subdirectories of the input folder to find all files. If the file has a `.py` extension, its file path and the file content will be stored in a dictionary. The dictionaries for all file are collected into a list. Reading all files into main memory increases the performance for downstream tasks, since main memory

³<https://github.com/d6t/d6tflow>



Figure 3.4: Schematic representation of the pipeline for RQ2 with configurable parameters. TODO describe

access is faster than disk access. Although the size of the systems main memory limits the dataset size, a file contains text and is therefore several kilobytes in size. The train dataset with 13,330 files is 86.47MB in accumulated file size.

Problem Detection

As a next step, the analysis plugins from the CCAP (see chapter 3.1.1) process every file and store the line number along with the problem type. As a result, for every file list of problematic line numbers and the corresponding type is available for further processing. The CCAP only covers the two analysis plugins for the problem types `CONDITION COMPARISON` and `RETURN NONE`. For model training, we introduce the simple algorithm to detect comparisons in conditions as problem type `CONDITION COMPARISON SIMPLE` (described in chapter 3.1.1)

Data Encoding

The data encoding step transform the internal data representation into input vectors x and output vectors y . The transformation is described in the following, multi-stage process.

Fixed Length Sequence The length of source code is dynamic, whereas the input size of our models is fixed. Therefore, the character stream of variable length has to be transformed into a token stream of fixed size. To extract meaningful tokens from the character stream, we use the python tokenizer from its standard library. The tokenizer separates the character stream based on the python syntax definition into tokens. All tokens contain a token type (like a name or operator token), the corresponding characters in the source code and a start and end position (line and column number). Next, the token stream is divided into a fixed size token sequence using a sliding window approach (size: 20, step size: 3).

Vocabulary Creation To represent a token value with a number, every token value needs a numeric representation. Therefore, the occurrence of each token value is counted and the index in an descending ordered list for each token value represents the token as an integer. To ignore potential capitalization mismatches, all token values are lower-case. The overall size of the vocabulary is configurable. If the size is smaller than the size of the distinct token value set, the least common token values will be replaced by an unknown token. The unknown token has a numeric representation one bigger than

the vocabulary size. We found that using a vocabulary size of 100,000 results in an acceptable 3.9% unknown token frequency.

Index-Based Encoding The textual token values are encoded with an Index-Based encoding. The most common token values have a numerical representation based on their index in a vocabulary (created with all token values of the train data). Optional, the type of each token is encoded by its numerical representation in the standard library. The type and value encoding are then combined alternating in a final encoding vector that represents an input sample.

Label Extraction With the internal representation, the ground truth is encoded as a line number. A transformation is necessary to label each sample corresponding to a given problem type. Label 1 is assigned, if the sample contains the given problem type. A sample contains a problem type, if it contains all tokens from the problematic line. If the sample is missing a token of the problematic line at the beginning or end, the label is 0 for non-problematic code.

Train/Test Split

TODO train test split before, or is it better here?

Code Manipulation

In research question TODO, we evaluate the models generalisation to detect similar pattern the model was not trained on. Therefore, we manipulate the code that it is still a rule violation, but the original analysis plugins would not spot them. The pipeline for code manipulation differ from the pipeline for RQ2 by an additional code manipulation task and a removal of the training step (see figure)

Return None For the problem type RETURN NONE, we manipulate the code to have an inline if statement that only returns None in one branch. Therefore, a function may still return None, but the analysis plugin would not detect this. Listing 3.4 shows an example for this case. Although a modification of the analysis plugin could cover the variations as well, we want to see how well machine learning models could perform on this task.




Figure 3.5: Pipeline for code manipulation and evaluation. TODO detailed description

To modify the original code, we use the preprocessed code with problems detected as in chapter 3.2.3. For every source code line containing the RETURN NONE problem, we use regular expressions to replace the *return None* with a variation as seen in 3.4.

The following data encoding step is similar to the one for model training, although this modified samples will only be used during evaluation.

```

def f(a,b):
    #detected by the analysis plugin
    return None

    #not detected by the analysis plugin
    return None if a < b else b

    #not detected by the analysis plugin
    return a if a < b else None
  
```

Listing 3.4: Samples for returning None. The first return would be flagged by the analysis plugin, the second and third return are modified variations that would be ignored by the analysis plugin. The performance of the machine learning models on detecting the latter will be evaluated.

Condition Comparison The analysis for the CONDITION COMPARISON problem type is an advancement of the analysis of the CONDITION COMPARISON SIMPLE

problem type. See listing 3.5 for examples of the different analysis results. Since the results of CONDITION COMPARISON are the modification of the CONDITION COMPARISON SIMPLE, we use the first as ground truth for evaluating the generalisation. Since the CONDITION COMPARISON SIMPLE problems are a subset of the CONDITION COMPARISON problems, we use a similar approach as before (see 3.2.3) to transform all problem occurrences to be undetectable by the CONDITION COMPARISON SIMPLE analysis plugin.

```
def f(a,b):
    #detected by CONDITION COMPARISON SIMPLE
    #detected by CONDITION COMPARISON
    if a < b:
        pass

    #not detected by CONDITION COMPARISON SIMPLE
    #detected by CONDITION COMPARISON
    if not a < b:
        pass

    #not detected by CONDITION COMPARISON SIMPLE
    #detected by CONDITION COMPARISON
    if isSmaller(a,b) or a < b:
        pass
```

Listing 3.5: Sample statements for the difference between the two analysis plugins CONDITION COMPARISON and CONDITION COMPARISON SIMPLE.

3.2.4 Models

Classifying code samples is a binary classification problem since we consider each problem type as a separate classification task. We train and evaluate a random forest classifier, a support-vector machine based classifier, a gradient boosting classifier and a neural network with LSTM cells.

Random Forest We use a random forest classifier with 100 decision trees to perform the binary classification. Additional to over- and undersampling, we experiment with an automatic class weighting inverse proportional to class frequency. The other hyperparameters follow the default implementation of the scikit-learn library: The quality measure of a split follows the gini function, the tree depth is unlimited and the number

of features to consider for each split is the square-root of the number of features overall.

Support Vector Machine We use the support vector classification from the scikit-learn library. As a kernel function, we choose the radial basis function for a performance baseline. Additionally, we try the class weighting inverse proportional to the class frequency to combat the imbalanced dataset.

Gradient Boosting Classifier For gradient boosting classifier, we mainly vary the number of boosting steps and the learning rate. Furthermore, we experiment with stochastic gradient boosting by setting the subsample parameter to 0.4 and 0.7. Other hyperparameter settings follow the default value of the scikit-learn library.

Neural Networks with LSTM Cells Our neural network consists of five layer (see listing 3.6). First, we use an embedding layer with an embedding size of 32. The input is the Index-Based encoded samples without type encoding. This layer is trained end-to-end with the complete network on the trainings-data. Furthermore, this layer covers most trainable parameters of the model. As a hidden layer, we use 10 LSTM cells with a prior and posterior dropout layer with a default 0.2 dropout. To perform binary classification, our model ends with a dense mapping onto a single neuron. The binary result is determined based on a 0.5 threshold for the output of the last neuron.

We vary the hyperparameters for batch size, number of epochs, embedding size and number of LSTM cells.

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 20, 32)	3200064

dropout (Dropout)	(None, 20, 32)	0

lstm (LSTM)	(None, 10)	1720

dropout_1 (Dropout)	(None, 10)	0

dense (Dense)	(None, 1)	11
=====		
Total params: 3,201,795		
Trainable params: 3,201,795		
Non-trainable params: 0		

Listing 3.6: Summary of our LSTM network. We use an embedding layer of size 32, 10 LSTM cells, a dropout layer before and after the LSTM layer with a default dropout ratio of 0.2 and a dense mapping to a single output neuron for binary classification with a threshold of 0.5.

4 Quantitative Evaluation

Research questions Description of the evaluation environment/setup Results/answers per research question Optional comparison to prior work Discussion and analysis of strengths/problems Note: see

4.1 Research Questions

RQ1 Can the Clean Code Analysis Platform be a useful addition to developers workflow besides existing tools?

RQ2 What models perform well on detecting non-clean code

4.1.1 RQ1: Can the Clean Code Analysis Platform be a useful addition to developers workflow besides existing tools?

Motivation Several tools are established that aim to improve code quality and detect unclean code. Every new tool with similar claims has to prove its usefulness beside the preexisting tools. We compare the CCAP with existing tools based on the design goals expandability, useability and integration. Additionally, we will mention and compare useful features unique to existing tools.

Approach To evaluate, if the CCAP is a useful addition, we compare different features to see, if it can supplement or replace existing tools. As described in section TODO, the compared tools are Sonarcube, PMD Source Code Analyser Project, Codacy and PyLint.

Results Table TODO shows a summary of the features.

Finding 1: Expandability For expandability of analysis plugins, CCAP is comparable to PyLint, that also offer plugins to analyse the raw string, the token stream or

the AST. PMD allows plugins to analyse the AST or define XPath rules. Since XPath rules can be used in the CCAP as well (using a third-party library in the plugin), PMD offers fewer expansion possibilities. Sonarqube, on the other hand, provides the most possibilities for extension. Not only can developers analyse the AST, specify XPath rules, but they have access to an additional semantic model of the code that provides direct access to structures such as methods, parameters, return values. This semantic model simplifies analysis in comparison to the AST analysis. Additionally, Sonarqube allows plugins to expose an API for other plugins to use. This different type of plugin further increases the possibilities for rule checking plugins. The least expandability offers Codacy, that allow customising or disabling existing rules, but does not offer to add rules to the existing ruleset.

Regarding output plugins, no compared tool offers the output customisation of CCAP. PMD and PyLint have different predefined output formats like JSON, HTML, CSV or text. Although PyLint allows customising the message format using a formatting string as a command-line argument, they do not offer extensions for the output. Sonarqube displays the analysis reports in its WebUI. The scanner tool, SonarScanner, sends the reports to the server component, that renders the results in the browser. Codacy's local scanner produces text output; the cloud version also has a WebUI.

Finding 2: Integration shortcomings The CCAP has significant shortcomings in its integrations into IDEs, build processes and CI pipelines. All contestant tools provide plugins for build systems like Gradle or Maven. Except for Codacy, all tools have IDE integrations into the most common editors. Sonarqube and Codacy have integrations into source control systems. Same applies to PMD and PyLint since they are included in Codacy and therefore have the same integrations.

Finding 3: Useability For a plugin developer, the useability of the plugin system is simple for CCAP and PyLint. Both tools have the same plugin mechanism and a simple plugin interface to work with. PMD has a similar, straightforward plugin interface, but its Java plugin has to be bundled into a jar file before adding to the classpath. The latter applies to Sonarqube as well; additionally, the powerful and rich plugin API makes it harder to use.

For the user, installing CCAP is like installing every other python package. Running from the command-line is fast and can be automated, e.g. using git pre-commit hooks. The same workflow would be achieved with PyLint. Sonarqube is more complex to

install due to its multiple components, but after the setup, the WebUI and integrations into most developer workflow have the best useability. Codacy as a cloud service only requires permission to access the source code repository to start scanning the code.

Finding 4: Multi-Language scanning Sonarqube, PMD and Codacy offer multi-language scanning. This multi-language ability is advantageous for code repositories with multiple languages and for reusing the same tools and infrastructure for multiple projects with different programming languages. Supporting multi-language scanning is not possible with CCAP and would require a redefinition of the architecture. Although multi-language support was not a design goal, it would be a significant advantage for adoption.

Finding 5: Maturity and Community-Support The maturity and community-support of tools impact the integration of the tools into the developer workflow. Some integrations are community-made and shared, so everybody has an advantage. The community support would also be necessary for CCAP to write and share additional analysis plugins and to integrate into different workflows.

Summary In summary, CCAP offers a good, but not best in class expandability with analysis plugins. It has a high useability for plugin developer and users. The lack of integrations for different workflows reflects the lack of community support and maturity. We see the CCAP as an addition to the powerful Sonarqube. The simple, yet robust expandability of the CCAP could supplement the shortcomings of Sonarqube in its powerful, but complex expandability. For instance, a code reviewer could code problematic code into a simple python plugin, distributes it to the team and would hopefully not reencounter the same problem. Additionally, the CCAP could be used to teach about Clean Code and to enforce specific coding rules to students, since the local setup is straightforward. Clean Code rules could be turned into analysis plugins by students or teachers without having to understand a complicated interface.

4.1.2 RQ2: What models perform well on detecting non-clean code?

Motivation As we have shown with the CCAP and the analysis plugins, it is possible to write an algorithm that can detect certain code patterns like non-clean code patterns. If we are able to write a well tested detection algorithm, we do not need error-prone

machine learning models to detect code patterns. Nonetheless, if we can not design an algorithm that detect a desired code pattern reliably, we may perform better using a machine learning model. Additionally, if we encode a subjective pattern, it may not be possible to design an algorithm to objectively detect such a pattern. As a solution, we could label code parts as our pattern and use machine learning approaches to implicitly extract rules to detect such a pattern.

For this research question, we will evaluate different machine learning models on different code patterns and compare the results.

Approach For answering the research question, we will use the approach detailed in chapter 3.2. The dataset, consisting of 18 projects from GitHub, is prepared as described in chapter 3.2.2. The data is split into 90% train and 10% test data. We use the train dataset to train Random Forest, Gradient Boosting Classifier, Support Vector Machines and a LSTM-based neural network. Due to the class imbalance, we will additional train with an oversampled dataset (oversampling rate of 0.5) and an undersampled dataset (undersampling rate of 0.5). Since our test set contains real-world code samples, we expect it to reperesent a real-world label distribution and we therefore accept potential effects of the data inbalance. Nevertheless, we will try to train with an over- and undersampled dataset to observe the effect on the performance.

Our evaluation metrics are recall, precision and the combined f1 score. Due to the data inbalance, accuracy is not a meaningful metric. A high recall means a high detection rate of non-clean code. The cost of a false negative prediction are potential costs in unclean code (maintainability, understandability). With a high precision, a detected non-clean code sample is likely to be a non-clean code sample and the system reports less false positives. The immediate cost of false positive predictions (resulting in a lower precision) are the extra developer time needed. Additionally, the cry wolf effect (TODO source) comes into play. In our case, a developer who has seen a false alarm will take subsequent alarms less serious. Additionally, a study has shown that a higher false alarm rate in advisory warning system in cars results in a lower, subjective evaluation of the system (TODO source Cooperative warning systems: The impact of false and unnecessary alarms on drivers' compliance). A low subjective evaluation of developers would lead to a decrease in adoption rate which consequently lead to more unfixed problematic code. Since both, recall and precision are important to the success of our models, we decided to use the equally weighted f1 score as single-value evaluation metric. Additionally, to compensate the cascading effect of a low precision, we require a precision

		RETURN_NONE	COND_COMP	COND_SIMPLE
Random Forest				
	F1	0.8328	0.9360	0.9375
	Recall	0.7179	0.9006	0.9004
	Precision	0.9913	0.9742	0.9779
SVM				
	F1			
	Recall			
	Precision			
Gradient Boosting Classifier				
	F1	0.4488	0.7958	0.7061
	Recall	0.9941	0.9727	0.9825
	Precision	0.2898	0.6734	0.5511
LSTM				
	F1	0.9927	0.9812	0.9825
	Recall	0.9912	0.9832	0.9875
	Precision	0.9941	0.9791	0.9776

Table 4.1: Best performing classifier for RQ2 (based on f1 score).

of 0.8 as a satisficing metric. This means, we accept all models that return two false positive and eight true positive samples. On the other hand, we define the recall as our optimizing metric.

We evaluate all models for the three different problem types RETURN_NONE, CONDITION_COMPARISON_SIMPLE and CONDITION_COMPARISON.

Results Table TODO shows the best performing hyperparameter per model for all three problem types. Figure TODO shows the confusion matrix those model configurations. We list all hyperparameter configurations and the corresponding training, train/test and test performance for all problem types in the appendix (Table A1-??).

Finding 1: poor SVM performance SVM performed poorly. Based on the f1 score, the best configuration has a f1 of 0.0066, a recall of 0.0134 and a precision of 0.0044. In other words, only 1.34% of positive samples are recognized as positive. The model seems to fail to separate the two classes. A reason may be the class imbalance, since only 0.21% of all samples are non-clean code (positive) with the RETURN_NONE type. In this configuration, we used undersampling with a ratio of 0.5, since oversampling was

not feasible due to the increase in train time (see section 3.2.1). After undersampling, 1/3 of all samples are therefore positive and 2/3 are negative. Since the time to train is still too long, we only took half of the undersampled dataset.

We additionally used the svm in a configuration with an automatic class weighing that is inverse-proportional to the class frequency. The recall score increases by 57-fold to 0.7660 whereas the precision decrease by half to 0.0026. We conclude, that the svm is very sensitive to class imbalance, which is an improper characteristic for our imbalanced classification problem. Additionally, a precision of 0.0026 is not useable for any application. We can not boost the precision with more data to help the model to distinguish true positives and false positives, since the train time will increase quadratically. We a baseline that low we do not expect much improvement with further investigation and therefore conclude that there is no valid reason to further invest in the svm model.

Finding 2: Random Forest Generally, the random forest classifier performs with a f1 score over 0.8 over all problem types with an oversampling ratio of 0.5, 100 trees, no additional class weighting and type encoding. For RN, this configuration has a 0.83 f1, 0.72 recall and 0.99 precision. For CCS and CC, the f1 with 0.94 are even better. The recall for those problem types increases to 0.9 while the precision remains high with 0.98 and 0.97 respectively. The random forest classifier therefore fullfills our satisficing condition of a precision higher than 0.8.

A common pattern is a worse performance on undersampled train sets. We observe a significant drop in performance for the RN problem type, if we use a 0.5 undersampling ratio. Similiarly, the performance for CCS and CC drops, although the decrease is smaller. Furthermore, we observe a moderate performance drop with 0.1 undersampling. From this pattern, we infer that a random forest classifier can profit from more data, even if they contain duplicated samples for the minority class due to oversampling. This can also explain the larger drop for undersampling for the RN problem type, since the non-clean code class contains less samples for RN than for CCS and CC (TODO numders of samples in train dataset). With less samples, the undersampling process reduces the overall amount of samples by deleting samples from the majority class. With a undersampling ratio of 0.1, less examples are deleted and the effect is less impactful, but observable.

The impact of class weighting in comparison to equivalent configurations seems to be neglectable. For oversampled datasets, the class weighting seems to slightly lower the performance whereas the opposite applies to undersampling. The former may be

explained by an overemphasise effect on the data. With oversampling, we change the the data distribution by duplicating the minority class and additionally emphasize the minority class with balancing. For the original dataset without under- or oversampling, class balancing does not have a significant impact on the performance for type encoding. For random forests without type encoding, we observe a slight 4 percentage points decrease in f1 performance.

With type encoding enabled, we see an improvement up to two percentage point in recall for all problem types with a comparable precision. Since the improvement is low and we observe a deterioration for undersampling, we can not infer a overall positive impact of type encoding on random forest. This requires additional investigations.

Overall, it seems like a random forest classifier can handle class inbalance fairly well, as long as distribution is not too unbalanced as for RN problem type. The performance deteriorates if the inbalance superceeds a certain threshold (for RN), otherwise it has minor impact. We base this conclusion on several observations: First, we see a better performance for the CCS and CC problem type overall, both of which have a more blanced distribution as the RN problem type. It is important to note, that the different problem types also have different difficulty levels for the model. But we assume the RN type is easier to detect for a model, since the model only has to learn two subsequent tokens anywhere in the sequence. Therefore, we attribute this effect to the class distribution, but we remark that this may be tested in an additional, isolated experiment. Second, without over- or undersampling, we observe similar performance for CCS and CC, despite their different frequency in the train set (TODO: class frequency values). Finally, we expect to have a bigger impact of class balancing, if the random forest classifier would be sensitive to class inbalance.

Finding 3: Gradient Boosting Classifier The overall best performing gradient boosting works with 200 boosting steps and a learning rate of 0.2. For the RN problem type, it delivers a f1 score of 0.8897, a recall score of 0.8257 and a precision of 0.9644. Compareable performance is measured for the CC type: 0.8665 f1, 0.8087 recall and 0.9332 precision. The performance decreases for the CCS type to 0.7738 f1, 0.687 recall and 0.8858 precision. This pattern applies to all configurations. Consequently, we assume it is harder for the gradient boosting classifier to learn the strucutre of the CCS problem type.

Using a subsampling of < 1.0 results in stochastic gradient boosting by using only a fraction of the samples for the base learners. Since we used this parameter in combination

with oversampling, the subsampled sample set contains all unique samples from the total train set. Therefore, the metrics do not differ from the non-subsampled counterpart.

TODO: subsample on non-oversampled dataset

For all over- and undersampling configurations, we can make two observations: First, the performance deteriorates compared to model trained without over- or undersampling. Second, we have better scores for the train datasets than for the train/test or test dataset (for all problem types). The latter is typical symptom of overfitting the train dataset and not generalizing well. Since we decrease the amount of train samples for undersampling and overemphasize samples of the minority class in oversampling, we decrease the variance in the train data and the model can not generalize very well. If we do not alter the label distribution in the train data, we observe for all problem types similar performance on the train and test set. The model does not overfit and generalize well to the test data. We conclude that the gradient boosting classifier is rather robust to class imbalance. TODO with more tests, review this - difference in train performance between oversampled, undersampled and not sampled \rightarrow profit from more samples, not sensitive to class imbalance. + we can increase the boosting step since it does not overfit the train data rather generalize very well (since test performance is same)

TODO: increase boost stages positive effect TODO: increased learn-rate positive?

5 Conclusion

An overall short summary of results What was successful, what not (and hypotheses why) Hints for further future work/extension

Bibliography

- [1] Go at Google: Language Design in the Service of Software Engineering.
- [2] Null Safety - Kotlin Programming Language.
- [3] Robert Baggen, José Pedro Correia, Katrin Schill, and Joost Visser. Standardized code quality benchmarking for improving software maintainability. 20(2):287–307.
- [4] Aurelien Delaitre, Bertrand Stivalet, Elizabeth Fong, and Vadim Okun. Evaluating Bug Finders – Test and Measurement of Static Code Analyzers. In *2015 IEEE/ACM 1st International Workshop on Complex Faults and Failures in Large Software Systems (COUFLESS)*, pages 14–20. IEEE.
- [5] Andrew Gerrand. Error handling and Go.
- [6] Tony Hoare. Null References: The Billion Dollar Mistake.
- [7] Robert C Martin. *Clean Code: A Handbook of Agile Software Craftsmanship*. Pearson Education.
- [8] Herbert Prahofer, Florian Angerer, Rudolf Ramler, and Friedrich Grillenberger. Static Code Analysis of IEC 61131-3 Programs: Comprehensive Tool Support and Experiences from Large-Scale Industrial Application. 13(1):37–47.
- [9] Mehwish Riaz, Emilia Mendes, and Ewan Tempero. A systematic review of software maintainability prediction and metrics. In *2009 3rd International Symposium on Empirical Software Engineering and Measurement*, pages 367–377. IEEE.
- [10] ISO Central Secretary. Systems and software engineering — Systems and software Quality Requirements and Evaluation (SQuaRE) — Guide to SQuaRE.
- [11] Dag I K Sjøberg, Dag Sjøberg, Bente Anda, and Audris Mockus. Questioning Software Maintenance Metrics: A Comparative Case Study. page 4.
- [12] Dave Wichers and Eitan Worcel. Source Code Analysis Tools | OWASP.

Parameter					RN			CCS			CC			
over-samp.	under-samp.	#trees	class weight	encode type		Train	Test	Holdout	Train	Test	Holdout	Train	Test	Holdout
1.0	-	100.0	None	True	<i>F1</i>	1.0	0.7036	0.8367	1.0	0.8886	0.9322	1.0	0.8852	0.9324
					<i>Rec1</i>	1.0	0.5432	0.7225	1.0	0.8103	0.8876	1.0	0.8116	0.8925
					<i>Prec</i>	1.0	0.9986	0.9937	1.0	0.9836	0.9815	1.0	0.9735	0.9761
0.5	-	100.0	balanced	True	<i>F1</i>	1.0	0.7	0.8385	1.0	0.8885	0.9321	1.0	0.8833	0.9329
					<i>Rec1</i>	1.0	0.5385	0.7249	1.0	0.8105	0.8882	1.0	0.8098	0.893
					<i>Prec</i>	1.0	1.0	0.9942	1.0	0.983	0.9806	1.0	0.9715	0.9765
0.5	-	100.0	None	True	<i>F1</i>	1.0	0.6973	0.8486	1.0	0.8887	0.9352	1.0	0.8837	0.9351
					<i>Rec1</i>	1.0	0.5353	0.7404	1.0	0.8105	0.8936	1.0	0.8084	0.8976
					<i>Prec</i>	1.0	1.0	0.9939	1.0	0.9837	0.9809	1.0	0.9744	0.9759
-	0.1	100.0	None	True	<i>F1</i>	1.0	0.6893	0.654	0.9999	0.8979	0.928	0.9999	0.8813	0.9345
					<i>Rec1</i>	1.0	0.8524	0.9275	1.0	0.8539	0.9243	0.9999	0.8101	0.9064
					<i>Prec</i>	1.0	0.5786	0.5051	0.9999	0.9466	0.9317	0.9999	0.9663	0.9644
-	-	100.0	balanced	True	<i>F1</i>	0.9999	0.5789	0.7503	0.9998	0.8494	0.9087	0.9998	0.8472	0.9121
					<i>Rec1</i>	1.0	0.4074	0.6018	1.0	0.7434	0.8404	1.0	0.7434	0.8484
					<i>Prec</i>	0.9997	1.0	0.9959	0.9996	0.9907	0.9891	0.9997	0.9847	0.9863
0.5	-	100.0	None	False	<i>F1</i>	1.0	0.5506	0.8519	1.0	0.7281	0.9217	1.0	0.721	0.9248
					<i>Rec1</i>	1.0	0.3799	0.7453	1.0	0.5781	0.8723	1.0	0.5733	0.8785
					<i>Prec</i>	1.0	1.0	0.9939	1.0	0.9833	0.977	1.0	0.9712	0.9764
1.0	-	100.0	None	False	<i>F1</i>	1.0	0.5229	0.8268	1.0	0.7511	0.9186	1.0	0.7322	0.9214
					<i>Rec1</i>	1.0	0.354	0.7084	1.0	0.6082	0.8674	1.0	0.5884	0.872
					<i>Prec</i>	1.0	1.0	0.9926	1.0	0.9819	0.9763	1.0	0.969	0.9767
-	-	100.0	None	True	<i>F1</i>	0.9999	0.3884	0.7666	0.9998	0.8513	0.9243	0.9998	0.8502	0.9269
					<i>Rec1</i>	0.9998	0.241	0.6226	0.9999	0.7446	0.8694	0.9998	0.7477	0.8759
					<i>Prec</i>	1.0	1.0	0.9972	0.9998	0.9938	0.9865	0.9998	0.9853	0.9841
-	-	100.0	balanced	False	<i>F1</i>	0.9999	0.3822	0.7405	0.9998	0.6828	0.8924	0.9998	0.6612	0.898
					<i>Rec1</i>	1.0	0.2363	0.5888	1.0	0.5207	0.8149	1.0	0.4976	0.8227
					<i>Prec</i>	0.9997	1.0	0.9976	0.9996	0.9914	0.9863	0.9997	0.985	0.9885
-	0.5	100.0	None	False	<i>F1</i>	1.0	0.3805	0.3024	1.0	0.7881	0.8157	1.0	0.8099	0.8624
					<i>Rec1</i>	1.0	0.9184	0.9712	1.0	0.9044	0.9684	1.0	0.8444	0.9566
					<i>Prec</i>	1.0	0.24	0.1791	1.0	0.6983	0.7046	1.0	0.7782	0.785
-	0.5	100.0	balanced	True	<i>F1</i>	1.0	0.2691	0.2357	1.0	0.7955	0.8181	1.0	0.849	0.871
					<i>Rec1</i>	1.0	0.9435	0.9775	1.0	0.949	0.9672	1.0	0.9254	0.9562
					<i>Prec</i>	1.0	0.157	0.134	1.0	0.6848	0.7088	0.9999	0.7842	0.7998
-	0.5	100.0	None	True	<i>F1</i>	1.0	0.2649	0.2227	1.0	0.7962	0.8151	1.0	0.8482	0.8693
					<i>Rec1</i>	1.0	0.9529	0.9782	1.0	0.9486	0.9709	1.0	0.9263	0.9602
					<i>Prec</i>	1.0	0.1538	0.1257	1.0	0.6859	0.7024	1.0	0.7823	0.7941
-	-	100.0	None	False	<i>F1</i>	0.9999	0.1415	0.7987	0.9998	0.6599	0.9106	0.9998	0.6359	0.9173
					<i>Rec1</i>	0.9998	0.0761	0.6658	0.9999	0.4936	0.8489	0.9998	0.4695	0.8571
					<i>Prec</i>	1.0	1.0	0.9979	0.9998	0.9951	0.982	0.9998	0.9847	0.9866

Table A1: random forest

Parameter						RN			CCS			CC			
over-samp.	under-samp.	stages	learn-rate	sub-sample	encode type		Train	Test	Holdout	Train	Test	Holdout	Train	Test	Holdout
-	-	300	0.2	1.0	True	<i>F1</i>	0.9364	0.7546	0.9328	0.8477	0.2405	0.8489	0.9087	0.1937	0.9088
						<i>Rec1</i>	0.9015	0.6444	0.9005	0.7926	0.1421	0.7953	0.8767	0.1168	0.8766
						<i>Prec</i>	0.9742	0.9102	0.9675	0.9111	0.782	0.9102	0.9431	0.5666	0.9434
-	-	200	0.2	1.0	True	<i>F1</i>	0.8554	0.6991	0.8499	0.7801	0.1906	0.7837	0.8644	0.1659	0.8667
						<i>Rec1</i>	0.7695	0.5691	0.765	0.6934	0.1092	0.6992	0.8057	0.097	0.8081
						<i>Prec</i>	0.9629	0.9062	0.956	0.8916	0.7501	0.8915	0.9324	0.5704	0.9344
-	-	200	0.2	0.4	True	<i>F1</i>	0.8554	0.6991	0.8499	0.7801	0.1906	0.7837	0.8644	0.1659	0.8667
						<i>Rec1</i>	0.7695	0.5691	0.765	0.6934	0.1092	0.6992	0.8057	0.097	0.8081
						<i>Prec</i>	0.9629	0.9062	0.956	0.8916	0.7501	0.8915	0.9324	0.5704	0.9344
-	-	200	0.2	0.7	True	<i>F1</i>	0.8554	0.6991	0.8499	0.7801	0.1906	0.7837	0.8644	0.1659	0.8667
						<i>Rec1</i>	0.7695	0.5691	0.765	0.6934	0.1092	0.6992	0.8057	0.097	0.8081
						<i>Prec</i>	0.9629	0.9062	0.956	0.8916	0.7501	0.8915	0.9324	0.5704	0.9344
-	-	100	0.2	1.0	True	<i>F1</i>	0.6438	0.5649	0.6441	0.6055	0.0999	0.6021	0.7417	0.1549	0.7437
						<i>Rec1</i>	0.4908	0.427	0.4896	0.4757	0.0541	0.472	0.6241	0.0884	0.6252
						<i>Prec</i>	0.9354	0.8344	0.9412	0.8329	0.652	0.8314	0.9138	0.6272	0.9175
0.5	-	200	0.2	1.0	True	<i>F1</i>	0.9936	0.5008	0.4165	0.9704	0.6022	0.7007	0.9624	0.517	0.791
						<i>Rec1</i>	0.9984	0.9498	0.9989	0.9843	0.6383	0.9842	0.9723	0.4903	0.9711
						<i>Prec</i>	0.9889	0.3401	0.2631	0.9568	0.5701	0.544	0.9528	0.5469	0.6672
-	0.5	200	0.2	1.0	True	<i>F1</i>	0.9944	0.4752	0.3817	0.9684	0.6098	0.6917	0.9632	0.5088	0.7905
						<i>Rec1</i>	0.9977	0.9796	0.9986	0.9816	0.6596	0.9817	0.9733	0.4861	0.972
						<i>Prec</i>	0.9912	0.3136	0.2359	0.9556	0.567	0.534	0.9533	0.5337	0.6661
0.5	-	100	0.2	1.0	True	<i>F1</i>	0.9831	0.3024	0.244	0.945	0.513	0.5866	0.9363	0.5504	0.7266
						<i>Rec1</i>	0.9914	0.9615	0.9954	0.9602	0.6923	0.962	0.938	0.5856	0.9383
						<i>Prec</i>	0.975	0.1794	0.1391	0.9302	0.4074	0.422	0.9347	0.5192	0.5928
-	0.5	100	0.2	1.0	True	<i>F1</i>	0.9843	0.2856	0.2278	0.943	0.5137	0.579	0.9381	0.5444	0.727
						<i>Rec1</i>	0.9935	0.9765	0.9958	0.957	0.7059	0.9588	0.9416	0.5789	0.9409
						<i>Prec</i>	0.9751	0.1673	0.1286	0.9294	0.4038	0.4147	0.9346	0.5138	0.5924
0.5	-	100	0.1	1.0	True	<i>F1</i>	0.9662	0.2085	0.1716	0.9111	0.4301	0.4863	0.8831	0.5131	0.613
						<i>Rec1</i>	0.9714	0.9482	0.9775	0.9217	0.7274	0.9242	0.868	0.6563	0.869
						<i>Prec</i>	0.9611	0.1171	0.0941	0.9007	0.3053	0.33	0.8988	0.4212	0.4735
-	0.5	100	0.1	1.0	True	<i>F1</i>	0.9662	0.1924	0.1586	0.9115	0.4394	0.4944	0.8853	0.5152	0.6216
						<i>Rec1</i>	0.9737	0.9537	0.9814	0.9194	0.7329	0.9215	0.8694	0.6492	0.872
						<i>Prec</i>	0.9587	0.107	0.0863	0.9038	0.3137	0.3379	0.9019	0.4271	0.4829

Table A2: gradient boosting classifier

Parameter								RN			CCS			CC		
over-samp.	under-samp.	emb-size	epochs	batch-size	#lstm-cells	dropout		Train	Test	Holdout	Train	Test	Holdout	Train	Test	Holdout
-	0.5	32	3	256	100.0	0.2 0.2	<i>F1</i> <i>Rec1</i> <i>Prec</i>	0.9994 1.0 0.9988	0.7719 0.9961 0.6301	0.818 1.0 0.6921	0.9986 0.9998 0.9974	0.2925 0.2295 0.4032	0.9548 0.9981 0.9152	0.998 0.9994 0.9966	0.2022 0.1273 0.4909	0.9722 0.9965 0.9491
-	-	32	2	512	10.0	0.2 0.2	<i>F1</i> <i>Rec1</i> <i>Prec</i>	0.9933 0.9916 0.9951	0.5996 0.4513 0.8929	0.9916 0.9919 0.9912	0.9905 0.9939 0.9872	0.164 0.0957 0.5713	0.9797 0.9818 0.9776	0.9868 0.996 0.9777	0.1886 0.1169 0.4877	0.977 0.99 0.9644
-	-	16	3	256	10.0	0.2 0.2	<i>F1</i> <i>Rec1</i> <i>Prec</i>	0.9899 0.998 0.982	0.3878 0.325 0.4808	0.9838 0.9961 0.9719	0.9908 0.9923 0.9893	0.1683 0.0975 0.6139	0.9812 0.9808 0.9815	0.9876 0.9879 0.9874	0.1785 0.1025 0.6921	0.976 0.975 0.9771
0.5	-	32	3	256	100.0	0.2 0.2	<i>F1</i> <i>Rec1</i> <i>Prec</i>	1.0 1.0 1.0	0.2892 0.3265 0.2595	0.9905 0.9905 0.9905	0.9997 1.0 0.9995	0.1579 0.0876 0.8024	0.9777 0.9732 0.9823	0.9993 1.0 0.9986	0.1673 0.0923 0.895	0.9756 0.9766 0.9747
-	-	32	3	512	10.0	0.2 0.2	<i>F1</i> <i>Rec1</i> <i>Prec</i>	0.9935 0.9889 0.9982	0.2594 0.1625 0.6429	0.9915 0.9887 0.9943	0.992 0.9972 0.9869	0.1667 0.094 0.7342	0.9809 0.988 0.974	0.9886 0.9967 0.9807	0.2444 0.153 0.6075	0.9774 0.9885 0.9665
-	-	32	3	64	10.0	0.2 0.2	<i>F1</i> <i>Rec1</i> <i>Prec</i>	0.9919 0.985 0.9988	0.1157 0.0644 0.5734	0.991 0.987 0.995	0.9897 0.9964 0.9831	0.1731 0.1103 0.4026	0.9808 0.9889 0.9728	0.9877 0.9919 0.9835	0.2024 0.119 0.6755	0.9781 0.9834 0.9728
-	-	32	3	256	10.0	0.2 0.2	<i>F1</i> <i>Rec1</i> <i>Prec</i>	0.9945 0.9936 0.9953	0.1135 0.1758 0.0838	0.9914 0.9947 0.9881	0.9928 0.992 0.9936	0.1698 0.0978 0.6438	0.9805 0.9779 0.9831	0.9906 0.9956 0.9857	0.2334 0.1435 0.6238	0.9783 0.9868 0.97
-	-	16	2	256	10.0	0.2 0.2	<i>F1</i> <i>Rec1</i> <i>Prec</i>	0.9925 0.9891 0.9958	0.065 0.0338 0.8776	0.9933 0.9933 0.9933	0.9864 0.9959 0.9771	0.1771 0.1024 0.6571	0.9791 0.9886 0.9698	0.9841 0.9763 0.992	0.1569 0.0869 0.8106	0.9734 0.9599 0.9872
-	-	32	2	256	10.0	0.2 0.2	<i>F1</i> <i>Rec1</i> <i>Prec</i>	0.9936 0.9912 0.996	0.0414 0.0306 0.0637	0.9933 0.9933 0.9933	0.9905 0.9952 0.9858	0.3125 0.2375 0.4567	0.981 0.9865 0.9755	0.9876 0.9941 0.9812	0.203 0.1205 0.6421	0.9778 0.9852 0.9706

Table A3: lstm